

Parallel algorithms

Introduction

Paolo Bientinesi

Umeå Universitet
pauldj@cs.umu.se

AQTIVATE workshop
28–29 November 2023



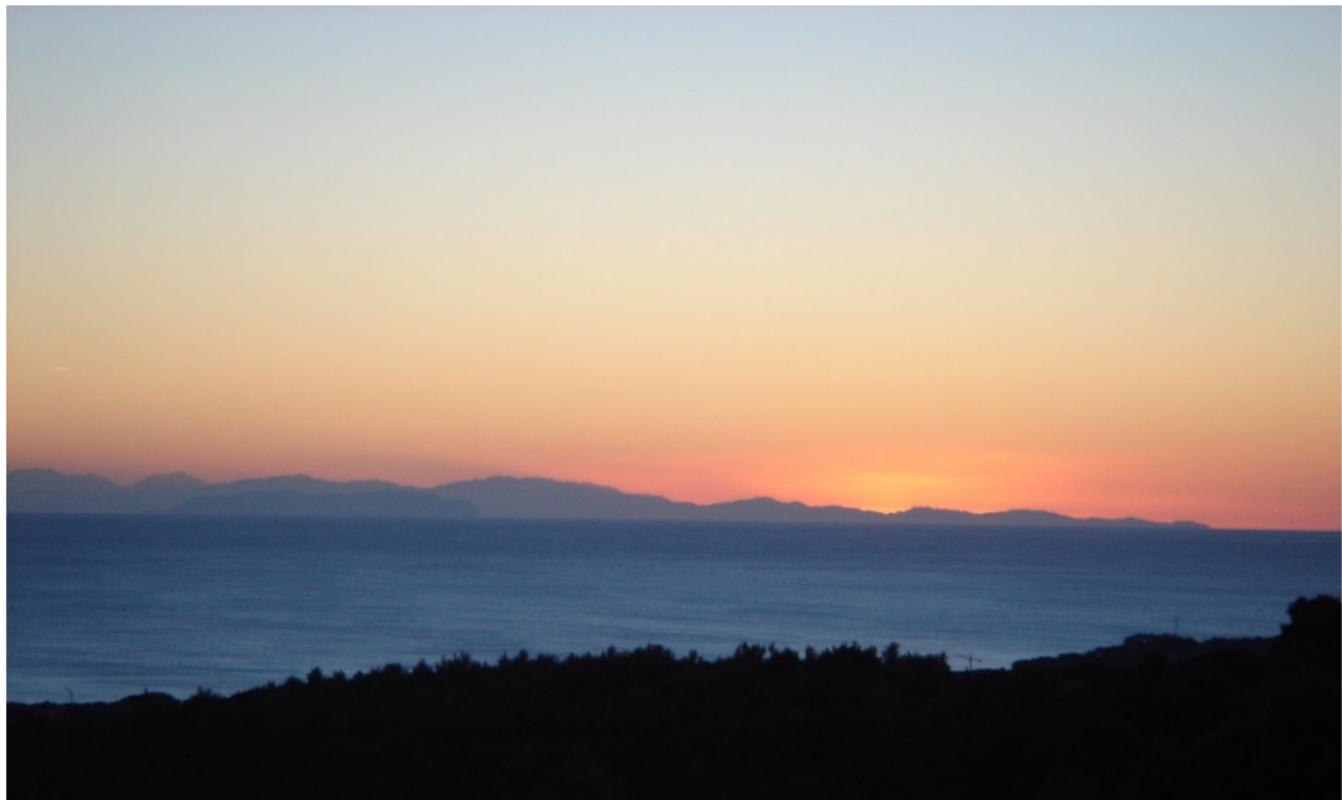
High Performance and
Automatic Computing

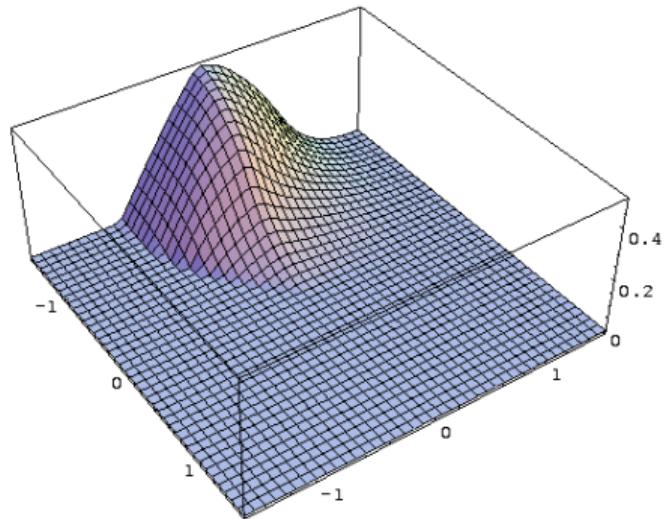
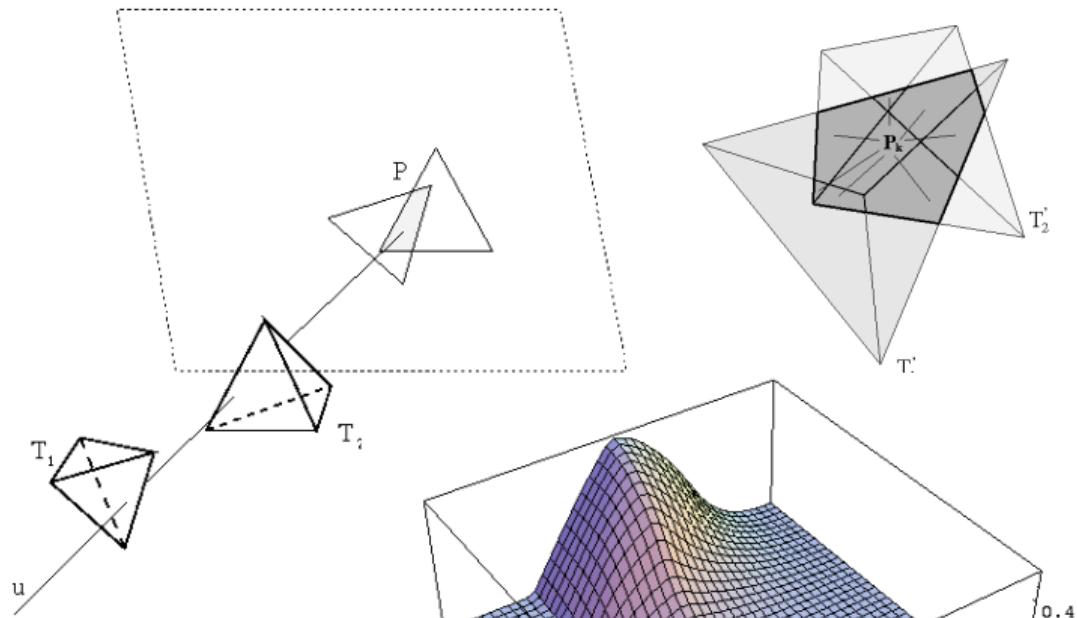


HPC2N

About me

Italy – Tuscany







- ▶ Symmetric eigenproblem
 $AX = X\Lambda$
- ▶ FLAME project
 Automatic generation of algorithms

Algorithm $LU = A$

Partition $A \rightarrow \left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right), L \rightarrow \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right), U \rightarrow \left(\begin{array}{c|c} U_{TL} & U_{TR} \\ \hline 0 & U_{BR} \end{array} \right)$
where A_{TL}, L_{TL}, U_{TL} , are 0×0

While $m(A_{TL}) \leq m(A)$ **do****Repartition**

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right), \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} L_{00} & 0 & 0 \\ \hline l_{10}^T & 1 & 0 \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right),$$

$$\left(\begin{array}{c|c} U_{TL} & U_{TR} \\ \hline 0 & U_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} U_{00} & u_{01} & U_{02} \\ \hline 0 & v_{11} & u_{12}^T \\ \hline 0 & 0 & U_{22} \end{array} \right)$$

where $\alpha_{11}, 1, v_{11}$ are scalars

$$v_{11} := \alpha_{11} - l_{10}^T u_{01}$$

$$u_{12}^T := a_{12}^T - l_{10}^T U_{02}$$

$$l_{21} := (a_{21} - L_{20} u_{01}) / v_{11}$$

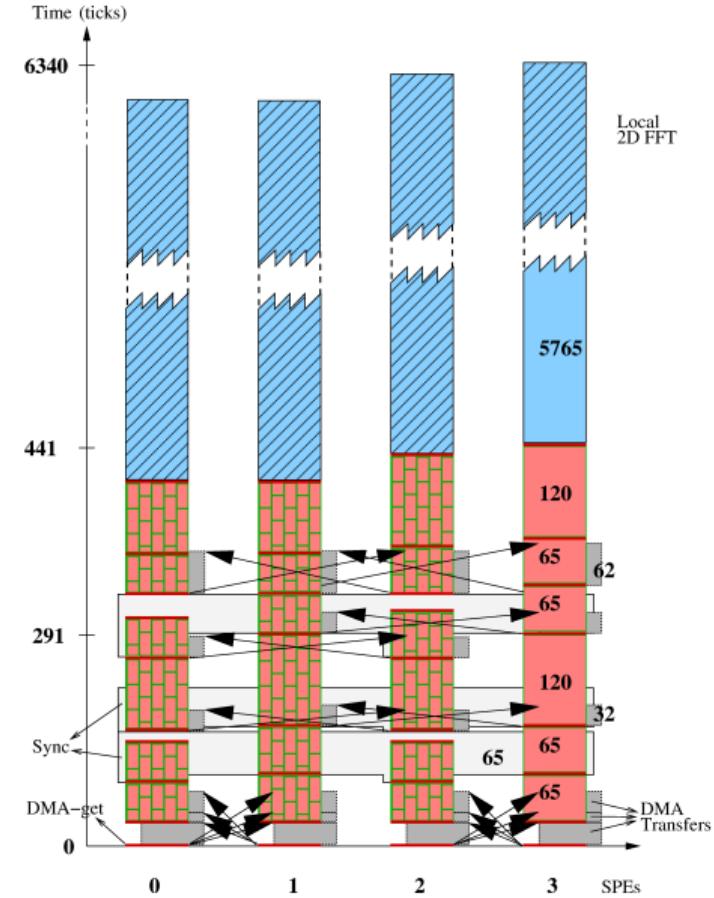
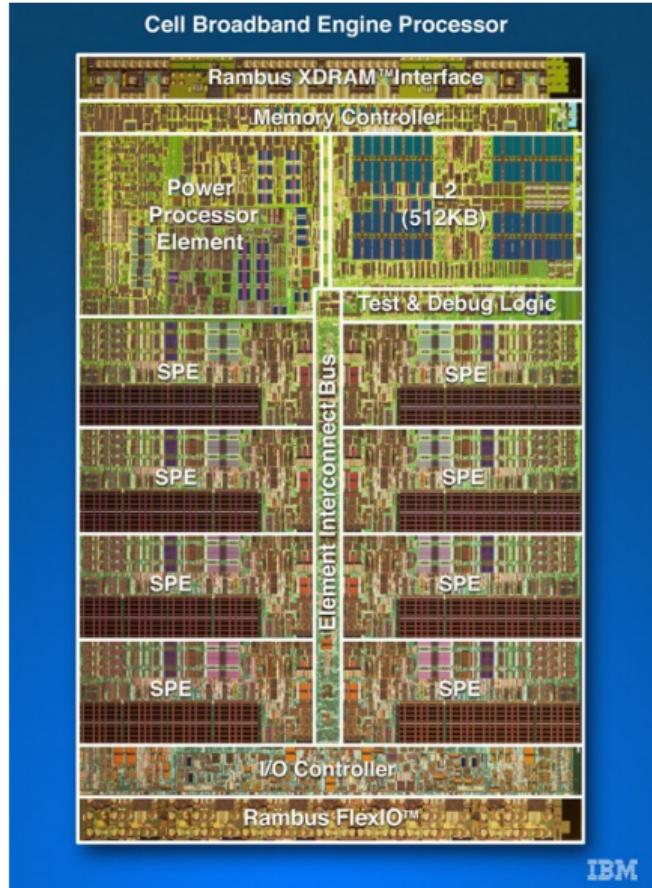
Continue with

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & a_{01} & A_{02} \\ \hline a_{10}^T & \alpha_{11} & a_{12}^T \\ \hline A_{20} & a_{21} & A_{22} \end{array} \right), \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} L_{00} & 0 & 0 \\ \hline l_{10}^T & 1 & 0 \\ \hline L_{20} & l_{21} & L_{22} \end{array} \right), \dots$$

endwhile



- ▶ Cell
- ▶ FFTs





High-Performance &
Automatic Computing





- ▶ Matrix & tensor operations
- ▶ HPC
- ▶ Computer music



High-Performance
Computing Center North

Why Parallel Computing?

- ▶ Time: Sequentially, the problem cannot be solved fast enough
 - ▶ Real-time constraints, accuracy, ...

Why Parallel Computing?

- ▶ Time: Sequentially, the problem cannot be solved fast enough
 - ▶ Real-time constraints, accuracy, ...
- ▶ Space: The problem cannot be solved on a single node
 - ▶ Large datasets, large memory requirement.

Why Parallel Computing?

- ▶ Time: Sequentially, the problem cannot be solved fast enough
 - ▶ Real-time constraints, accuracy, ...
- ▶ Space: The problem cannot be solved on a single node
 - ▶ Large datasets, large memory requirement.
- ▶ My computer is parallel, why not to take advantage of it?

- ▶ List of the 500 fastest supercomputers in the world
- ▶ Twice per year (ISC, June, Germany. SC, November, USA)
- ▶ Computers ranked based on the LINPACK benchmark
 - ▶ Solution of linear system of equations: $Ax = b$
 - ▶ Result measured in FLOPS/s¹ (in double precision)
- ▶ Established in 1993: 60 GFlop/s
- ▶ June 2016: 122,000,000 GFlop/s
- ▶ June 2023: 1,194,000,000 GFlop/s

¹Floating Point Operations per Second

Examples of supercomputers



Fugaku

Rank #1: June 2020 – Nov. 2021

Source: top500.org

Fugaku

- ▶ Site: Riken Center for Computational Science in Kobe, Japan
- ▶ 158,976 nodes
 - ▶ Fujitsu A64FX CPU (48+4 core)
- ▶ 7,630,848 cores
- ▶ 32 GiB/node RAM
- ▶ 442,010 PFlop/s (537,212 PFlop/s)
- ▶ 29,899 kW

Examples of supercomputers



Sunway TaihuLight

Rank #1: June 2017 – Nov. 2017

Source: top500.org

Sunway TaihuLight

- ▶ Site: National Supercomputing Center in Wuxi, China
- ▶ 40,960 SW26010 processors, each
 - ▶ 256 processing cores
 - ▶ 4 auxiliary cores
- ▶ 10,649,600 cores
- ▶ 1,310,720 GB RAM
- ▶ 93,014.6 TFlop/s (125,435.9 TFlop/s)
- ▶ 15,371 kW

So...

is parallel programming only for supercomputers?

So...

is parallel programming only for supercomputers?

Not at all!!

Parallel Computers are everywhere!



Source: hp.com



Source: indiatimes.com



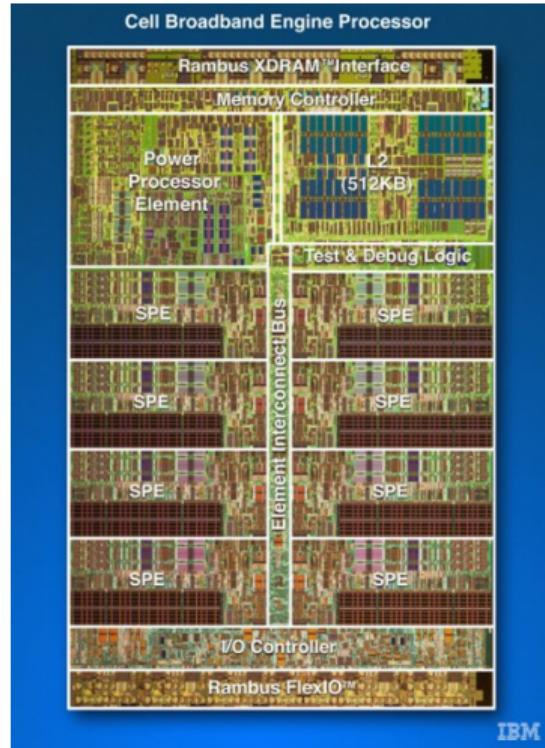
Source: bq.com



Source: wearabledevices.es

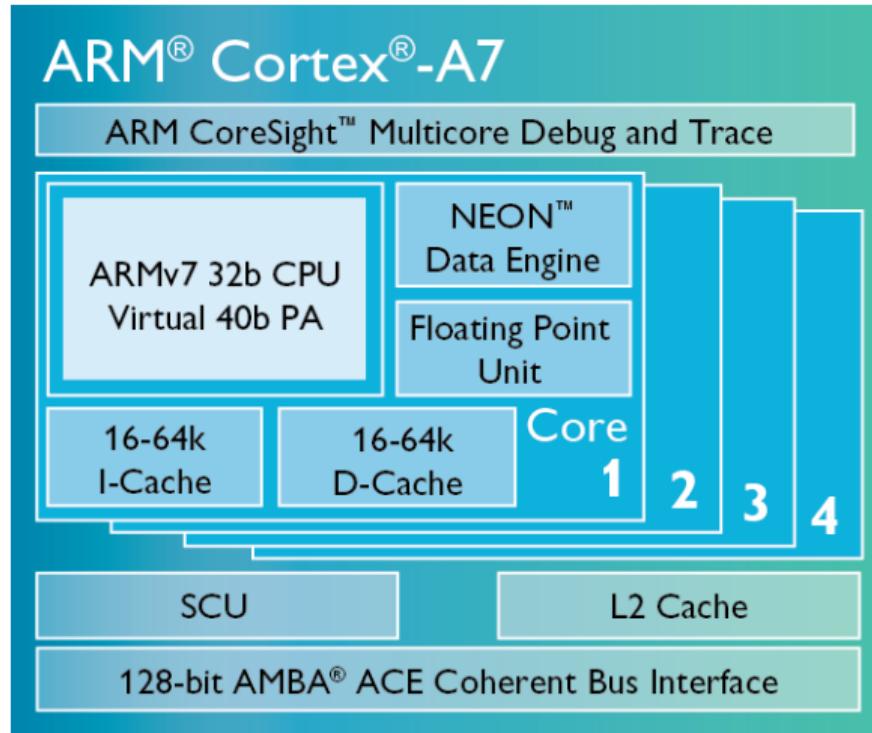
Parallel Computers are everywhere!

Play Station 3



Parallel Computers are everywhere!

Cell phones



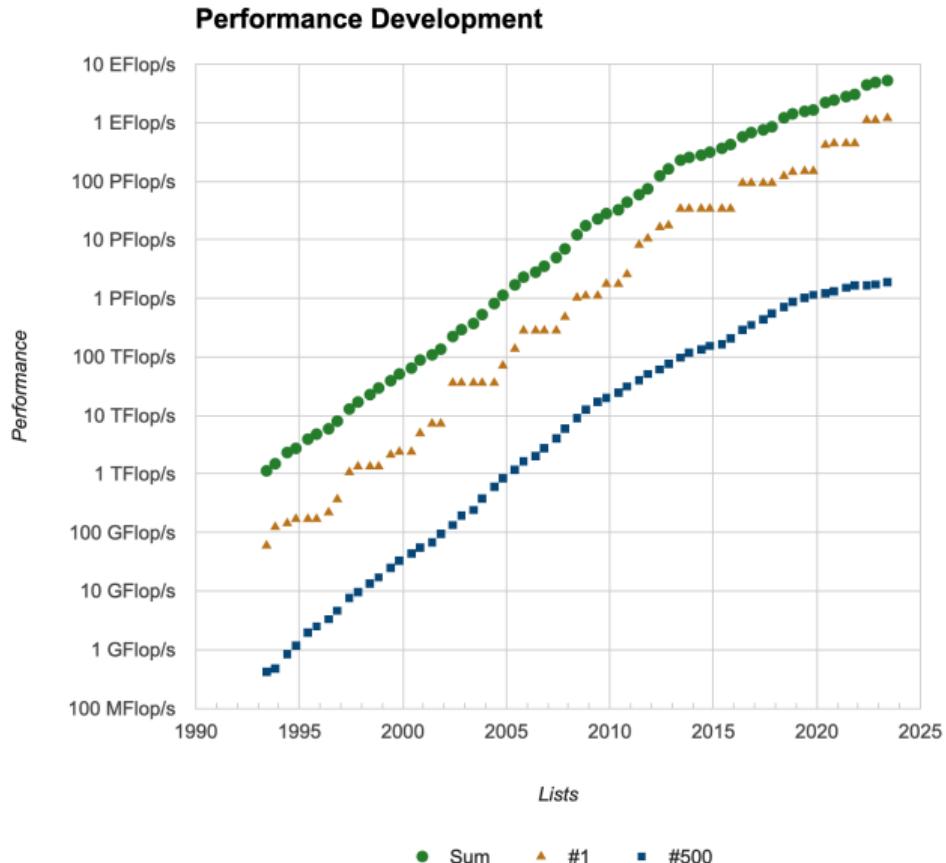
Source: arm.com

Fundamental axiom of High-Performance Computing

Fundamental axiom of High-Performance Computing

Every computer is a supercomputer, even your laptop!

Top 500 supercomputers in the world



My laptop > #1 supercomputer of 1997 !!!



Analogy: Supercars

We all own Formula-1 racing cars



Analogy: Supercars

We all own Formula-1 racing cars



do we drive them as such?

Analogy: highly parallel computers



...often used sequentially



Sequential pseudo-code

Sequential pseudo-code

Dijkstra shortest path (from Wikipedia)

```
1  function Dijkstra(GRAPH, SOURCE):
2
3      create vertex set Q
4
5      for each vertex v in GRAPH:
6          dist[v] ← INFINITY
7          prev[v] ← UNDEFINED
8          add v to Q
9          dist[SOURCE] ← 0
10
11     while Q is not empty:
12         u ← vertex in Q with min dist[u]
13
14         remove u from Q
15
16         for each neighbor v of u:           // only v that are still in Q
17             alt ← dist[u] + length(u, v)
18             if alt < dist[v]:
19                 dist[v] ← alt
20                 prev[v] ← u
21
22     return dist[], prev[]
```

Sequential pseudo-code

Dijkstra shortest path (from Wikipedia)

```
1  function Dijkstra(GRAPH, SOURCE):
2
3      create vertex set Q
4
5      for each vertex v in GRAPH:
6          dist[v] ← INFINITY
7          prev[v] ← UNDEFINED
8          add v to Q
9          dist[SOURCE] ← 0
10
11     while Q is not empty:
12         u ← vertex in Q with min dist[u]
13
14         remove u from Q
15
16         for each neighbor v of u:           // only v that are still in Q
17             alt ← dist[u] + length(u, v)
18             if alt < dist[v]:
19                 dist[v] ← alt
20                 prev[v] ← u
21
22     return dist[], prev[]
```

Top-to-bottom execution: **One line of the program is completed before the next one is executed.**

Sequential code

Dijkstra shortest path (from [geeksforgeeks.org](https://www.geeksforgeeks.org/dijkstras-shortest-path-algorithm-greedy-method/))

```
#include <limits.h>
#include <stdio.h>
#include <stdbool.h>
int minDistance(int dist[], bool sptSet[]);
void printSolution(int dist[]);

void dijkstra(int graph[V][V], int src) {
    int dist[V];
    bool sptSet[V];

    for (int i = 0; i < V; i++)
        dist[i] = INT_MAX, sptSet[i] = false;
    dist[src] = 0;

    // Find shortest path for all vertices
    for (int count = 0; count < V - 1; count++) {
        int u = minDistance(dist, sptSet);
        sptSet[u] = true;

        for (int v = 0; v < V; v++)
            if (!sptSet[v] && graph[u][v] && dist[u] != INT_MAX
                && dist[u] + graph[u][v] < dist[v])
                dist[v] = dist[u] + graph[u][v];
    }
    printSolution(dist);
}
```

Sequential code

Dijkstra shortest path (from [geeksforgeeks.org](https://www.geeksforgeeks.org/dijkstras-shortest-path-algorithm-greedy-method/))

```
#include <limits.h>
#include <stdio.h>
#include <stdbool.h>
int minDistance(int dist[], bool sptSet[]);
void printSolution(int dist[]);

void dijkstra(int graph[V][V], int src) {
    int dist[V];
    bool sptSet[V];

    for (int i = 0; i < V; i++)
        dist[i] = INT_MAX, sptSet[i] = false;
    dist[src] = 0;

    // Find shortest path for all vertices
    for (int count = 0; count < V - 1; count++) {
        int u = minDistance(dist, sptSet);
        sptSet[u] = true;

        for (int v = 0; v < V; v++)
            if (!sptSet[v] && graph[u][v] && dist[u] != INT_MAX
                && dist[u] + graph[u][v] < dist[v])
                dist[v] = dist[u] + graph[u][v];
    }
    printSolution(dist);
}
```

The code dictates how the computational steps unfold. Everything else is left to the compiler.

Parallel code

$$v_i^{\text{next}} := f(v_{i-1}^{\text{now}}) + f(v_i^{\text{now}}) - v_{i+1}^{\text{now}}$$

Parallel code

$$v_i^{\text{next}} := f(v_{i-1}^{\text{now}}) + f(v_i^{\text{now}}) - v_{i+1}^{\text{now}}$$

```
MPI_Isend(v_now, n, MPI_INT, proc_left, proc_left, MPI_COMM_WORLD, &req_send_left); // v_i -> p_{i-1}
MPI_Irecv(v_left, n, MPI_INT, proc_left, proc_left, MPI_COMM_WORLD, &req_recv_left); // f(v_{i-1}) <- p_{i-1}
MPI_Irecv(v_right, n, MPI_INT, proc_right, me, MPI_COMM_WORLD, &req_recv_right); // v_{i+1} <- p_{i+1}

f(v_now, v_next, n); // v_next = f(v_now)

// copy fv and send to next
memcpy(fv, v_next, n * sizeof(int));
MPI_Isend(fv, n, MPI_INT, proc_right, me, MPI_COMM_WORLD, &req_send_right);

// Receive in any order
MPI_Request reqs2 = {req_recv_left, req_recv_right};
int idx, p;

for (p=0; p<2; p++) {
    MPI_Waitany(2, reqs, &idx, MPI_STATUS_IGNORE);
    if (idx == 0) // Received from proc on the left
    {
        printf("%d Received from left (p%d): %d\n", me, proc_left, v_left0);
        for (i = 0; i < n; i++)
            v_nexti += v_lefti;
    } else // Received from proc on the right
    {
        printf("%d Received from right (p%d): %d\n", me, proc_right, v_right0);
        for (i = 0; i < n; i++)
            v_nexti -= v_righti;
    }
}
// Wait for the sends to complete
MPI_Wait(&req_send_left, MPI_STATUS_IGNORE);
MPI_Wait(&req_send_right, MPI_STATUS_IGNORE);
```

RED: computation BLACK: synchronization, data movement

No assumptions can be made about which processor starts or completes first.

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?

Not really. Maybe in 5-10 years?

(maybe, but do not believe the hype)

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?
Not really. Maybe in 5-10 years? (maybe, but do not believe the hype)
- ▶ Manually?

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?

Not really. Maybe in 5-10 years?

(maybe, but do not believe the hype)

- ▶ Manually? Yes, very much so.

With (a slew of) tools, libraries, programming languages for parallel computing.

E.g.: Charm, Chapel, Cilk, Coarray, CUDA, Hadoop, MapReduce, OpenCL, OpenSHMEM, OpenACC, Pthreads, TBB, ... (and many more)

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?

Not really. Maybe in 5-10 years?

(maybe, but do not believe the hype)

- ▶ Manually? Yes, very much so.

With (a slew of) tools, libraries, programming languages for parallel computing.

E.g.: Charm, Chapel, Cilk, Coarray, CUDA, Hadoop, MapReduce, OpenCL, OpenSHMEM, OpenACC, Pthreads, TBB, ... (and many more)

- ▶ Unfortunately, too many, too different, too differently supported, ...
- ▶ Difficult to discern research projects from reliable long-term efforts
- ▶ Difficult to predict support and involvement from all the different stakeholders

Holy grail: Automatic parallelization

How to decompose? How to coordinate?

- ▶ Automatic languages, parallelizing compilers?

Not really. Maybe in 5-10 years?

(maybe, but do not believe the hype)

- ▶ Manually? Yes, very much so.

With (a slew of) tools, libraries, programming languages for parallel computing.

E.g.: Charm, Chapel, Cilk, Coarray, CUDA, Hadoop, MapReduce, OpenCL, OpenSHMEM, OpenACC, Pthreads, TBB, ... (and many more)

- ▶ Unfortunately, too many, too different, too differently supported, ...
- ▶ Difficult to discern research projects from reliable long-term efforts
- ▶ Difficult to predict support and involvement from all the different stakeholders
- ▶ Best bet: Invest in those core tools/libraries that existed already 10 years ago, and are still being actively developed: MPI, pthreads, OpenMP, CUDA

Summarizing

- ▶ Parallel programming is critical in science and engineering

Summarizing

- ▶ Parallel programming is critical in science and engineering
- ▶ Relevant not only to supercomputers, but to any workstation/laptop

Summarizing

- ▶ Parallel programming is critical in science and engineering
- ▶ Relevant not only to supercomputers, but to any workstation/laptop
- ▶ Different platforms exhibit different architectural features, and different types and degrees of parallelism.

Summarizing

- ▶ Parallel programming is critical in science and engineering
- ▶ Relevant not only to supercomputers, but to any workstation/laptop
- ▶ Different platforms exhibit different architectural features, and different types and degrees of parallelism.
- ▶ Different forms of parallelism are handled via different programming paradigms:
Vectorization – Data Parallelism – Multi-threading – Message Passing

Summarizing

- ▶ Parallel programming is critical in science and engineering
- ▶ Relevant not only to supercomputers, but to any workstation/laptop
- ▶ Different platforms exhibit different architectural features, and different types and degrees of parallelism.
- ▶ Different forms of parallelism are handled via different programming paradigms:
 Vectorization – Data Parallelism – Multi-threading – Message Passing
- ▶ More often than not, the best algorithmic solution depends on the specific paradigm.

Summarizing

- ▶ Parallel programming is critical in science and engineering
- ▶ Relevant not only to supercomputers, but to any workstation/laptop
- ▶ Different platforms exhibit different architectural features, and different types and degrees of parallelism.
- ▶ Different forms of parallelism are handled via different programming paradigms:
 Vectorization – Data Parallelism – Multi-threading – Message Passing
- ▶ More often than not, the best algorithmic solution depends on the specific paradigm.
- ▶ In a nutshell: **The burden is and will be on the programmer.** (Hence workshops like this one)

Vectorization

- ▶ A processor operates on data that resides in the registers.

Vectorization

- ▶ A processor operates on data that resides in the registers.
- ▶ Nowadays, one single register contains 2–16 scalars; the exact amount depends on the register width (128, 256, 512bits) and on the data type (single, double, ...).

Vectorization

- ▶ A processor operates on data that resides in the registers.
- ▶ Nowadays, one single register contains 2–16 scalars; the exact amount depends on the register width (128, 256, 512bits) and on the data type (single, double, ...).
- ▶ Hence, the processor operates on vectors!

Vectorization

- ▶ A processor operates on data that resides in the registers.
- ▶ Nowadays, one single register contains 2–16 scalars; the exact amount depends on the register width (128, 256, 512bits) and on the data type (single, double, ...).
- ▶ Hence, the processor operates on vectors!
- ▶ Let α_i , β_i and γ_i be scalars (for all i 's).
Even if you want to only perform $\gamma_1 := \alpha_1 + \beta_1$, the processor is equipped to simultaneously execute
$$\{\gamma_1 := \alpha_1 + \beta_1, \quad \gamma_2 := \alpha_2 + \beta_2, \quad \dots, \quad \gamma_w := \alpha_w + \beta_w\}$$

w depend on the register width and the datatype of the scalars.

Vectorization

- ▶ A processor operates on data that resides in the registers.
- ▶ Nowadays, one single register contains 2–16 scalars; the exact amount depends on the register width (128, 256, 512bits) and on the data type (single, double, ...).
- ▶ Hence, the processor operates on vectors!
- ▶ Let α_i , β_i and γ_i be scalars (for all i 's).
Even if you want to only perform $\gamma_1 := \alpha_1 + \beta_1$, the processor is equipped to simultaneously execute
$$\{\gamma_1 := \alpha_1 + \beta_1, \quad \gamma_2 := \alpha_2 + \beta_2, \quad \dots, \quad \gamma_w := \alpha_w + \beta_w\}$$

w depend on the register width and the datatype of the scalars.
- ▶ Maximum possible speedup: w .

SIMD paradigm

- ▶ SIMD = Single Instruction Multiple Data

SIMD paradigm

- ▶ SIMD = Single Instruction Multiple Data
- ▶ By hand: Streaming SIMD Extensions. Almost like Assembly.
“MMX”, “SSE”, “SSE2”, “SSE3”, “SSE4”, “AVX”, “AVX2”, “AVX-512”, “NEON”, ...

SIMD paradigm

- ▶ SIMD = Single Instruction Multiple Data
- ▶ By hand: Streaming SIMD Extensions. Almost like Assembly.
“MMX”, “SSE”, “SSE2”, “SSE3”, “SSE4”, “AVX”, “AVX2”, “AVX-512”, “NEON”, ...
- ▶ Autovectorization.

SIMD paradigm

- ▶ SIMD = Single Instruction Multiple Data
- ▶ By hand: Streaming SIMD Extensions. Almost like Assembly.
“MMX”, “SSE”, “SSE2”, “SSE3”, “SSE4”, “AVX”, “AVX2”, “AVX-512”, “NEON”, ...
- ▶ Autovectorization.
- ▶ Via directives: OpenMP.

Data-parallelism

- ▶ Closely related to vectorization: coarser granularity.

Data-parallelism

- ▶ Closely related to vectorization: coarser granularity.
- ▶ w might be even in the 100s and 1000s.
Large speedups expected wrt sequential execution.

Data-parallelism

- ▶ Closely related to vectorization: coarser granularity.
- ▶ w might be even in the 100s and 1000s.
Large speedups expected wrt sequential execution.
- ▶ Typical in accelerators and “many-core” architectures.

Data-parallelism

- ▶ Closely related to vectorization: coarser granularity.
- ▶ w might be even in the 100s and 1000s.
Large speedups expected wrt sequential execution.
- ▶ Typical in accelerators and “many-core” architectures.
- ▶ Identify program flow that is identical for different data.

Data-parallelism

- ▶ Closely related to vectorization: coarser granularity.
- ▶ w might be even in the 100s and 1000s.
Large speedups expected wrt sequential execution.
- ▶ Typical in accelerators and “many-core” architectures.
- ▶ Identify program flow that is identical for different data.
- ▶ OpenCL, Cuda (GPUs).

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.
- ▶ Feature: Main memory is shared among threads.
Some of the cache memories might also be shared.

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.
- ▶ Feature: Main memory is shared among threads.
Some of the cache memories might also be shared.
- ▶ No need for communication, BUT, need for synchronization.

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.
- ▶ Feature: Main memory is shared among threads.
Some of the cache memories might also be shared.
- ▶ No need for communication, BUT, need for synchronization.
- ▶ Example: students working on a large calculation at the blackboard.

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.
- ▶ Feature: Main memory is shared among threads.
Some of the cache memories might also be shared.
- ▶ No need for communication, BUT, need for synchronization.
- ▶ Example: students working on a large calculation at the blackboard.
- ▶ pthreads, OpenMP

Shared-memory parallelism – Multi-threading

- ▶ Multiple cores, multiple threads per core.
- ▶ Feature: Main memory is shared among threads.
Some of the cache memories might also be shared.
- ▶ No need for communication, BUT, need for synchronization.
- ▶ Example: students working on a large calculation at the blackboard.
- ▶ pthreads, OpenMP
- ▶ Max speedup: Typically in the 10s.

Multi-threading

- ▶ Iteration-based parallelism.

Multi-threading

- ▶ Iteration-based parallelism.
- ▶ Task-based parallelism.

Multi-threading

- ▶ Iteration-based parallelism.
- ▶ Task-based parallelism.
- ▶ Dependencies.

Multi-threading

- ▶ Iteration-based parallelism.
- ▶ Task-based parallelism.
- ▶ Dependencies.
- ▶ Scheduling; pinning; false-sharing.

Multi-threading

- ▶ Iteration-based parallelism.
- ▶ Task-based parallelism.
- ▶ Dependencies.
- ▶ Scheduling; pinning; false-sharing.
- ▶ Synchronization; race-conditions, mutual exclusion.

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.
- ▶ Data distributed among participants.

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.
- ▶ Data distributed among participants.
- ▶ Feature: Independent computing.
Data has to be exchanged = communication.
It is possible to overlap communication with computation.

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.
- ▶ Data distributed among participants.
- ▶ Feature: Independent computing.
Data has to be exchanged = communication.
It is possible to overlap communication with computation.
- ▶ Example: students working independently on (parts of) the same homework, and exchanging information on the phone.

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.
- ▶ Data distributed among participants.
- ▶ Feature: Independent computing.
Data has to be exchanged = communication.
It is possible to overlap communication with computation.
- ▶ Example: students working independently on (parts of) the same homework, and exchanging information on the phone.
- ▶ MPI: Message Passing Interface

Distributed-memory parallelism – Message passing

- ▶ Many nodes, each with its own local memory.
From few to 10^5+ nodes.
- ▶ Data distributed among participants.
- ▶ Feature: Independent computing.
Data has to be exchanged = communication.
It is possible to overlap communication with computation.
- ▶ Example: students working independently on (parts of) the same homework, and exchanging information on the phone.
- ▶ MPI: Message Passing Interface
- ▶ Max speedup: Unbounded.

Message passing

- ▶ Data distribution

Message passing

- ▶ Data distribution
- ▶ Communication primitives.

Message passing

- ▶ Data distribution
- ▶ Communication primitives.
- ▶ Collective communication.

Message passing

- ▶ Data distribution
- ▶ Communication primitives.
- ▶ Collective communication.
- ▶ Point-to-point communication.

Message passing

- ▶ Data distribution
- ▶ Communication primitives.
- ▶ Collective communication.
- ▶ Point-to-point communication.
- ▶ One-sided communication.

Message passing

- ▶ Data distribution
- ▶ Communication primitives.
- ▶ Collective communication.
- ▶ Point-to-point communication.
- ▶ One-sided communication.
- ▶ Synchronous & asynchronous communication.

General pattern for parallel algorithms

1. Decomposition of the problem into sub-problems
2. Parallel solution of the sub-problems
3. Composition of the sub-solutions