

Performance Modeling

Analytic performance modeling and its use
in scientific computing

Georg Hager, NHR@FAU

R.W. Hockney and I.J. Curington: $f_{1/2}$: A parameter to characterize memory and communication bottlenecks. Parallel Computing 10, 277-286 (1989). DOI: [10.1016/0167-8191\(89\)90100-2](https://doi.org/10.1016/0167-8191(89)90100-2)

W. Schönauer: Scientific Supercomputing: Architecture and Use of Shared and Distributed Memory Parallel Computers. Self-edition (2000)

S. Williams: Auto-tuning Performance on Multicore Computers. UCB Technical Report No. UCB/EECS-2008-164. PhD thesis (2008)



Motivation

- Analytic performance modeling:

An analytic white-box performance model is a **simplified mathematical description** of the **hardware** and its interaction with **software**. It is able to **predict the runtime/performance** of code from “**first principles**.”

- Basic questions addressed by **analytic performance models**

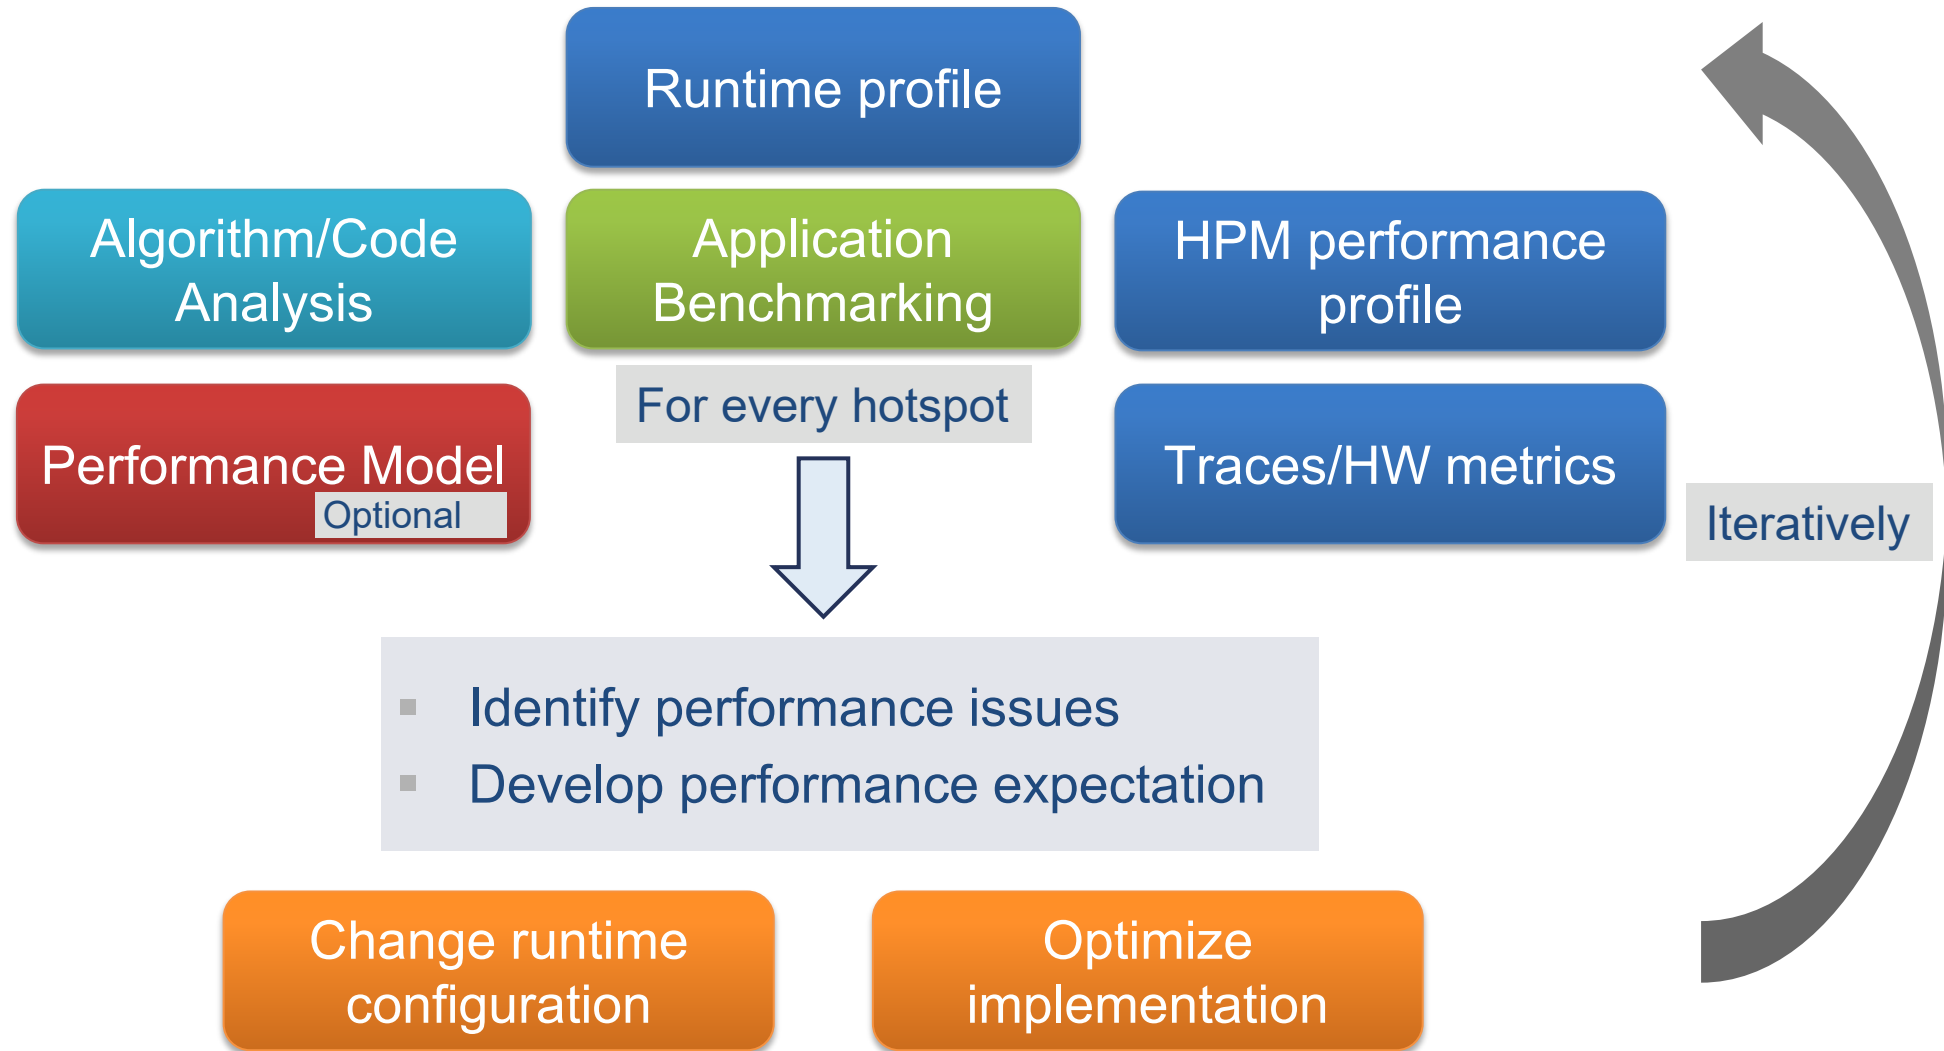
- What is the bottleneck?
- What is the next bottleneck after optimization?
- Impact of hardware features → **co-design, architectural exploration**

- What if the **model fails**?

- We **learn something**
- We may still be able to use the model in **a less predictive way**

Performance
Engineering (PE)

Performance Engineering is a process



Getting a little more general

What data/knowledge can a model be based on?

- Only documented hardware properties + hypotheses

- Purely analytic model

} white
box

- Hardware properties + measurements + hypotheses

- (Partly) phenomenological model


} gray
box

- Measured performance/speedup data + hypotheses

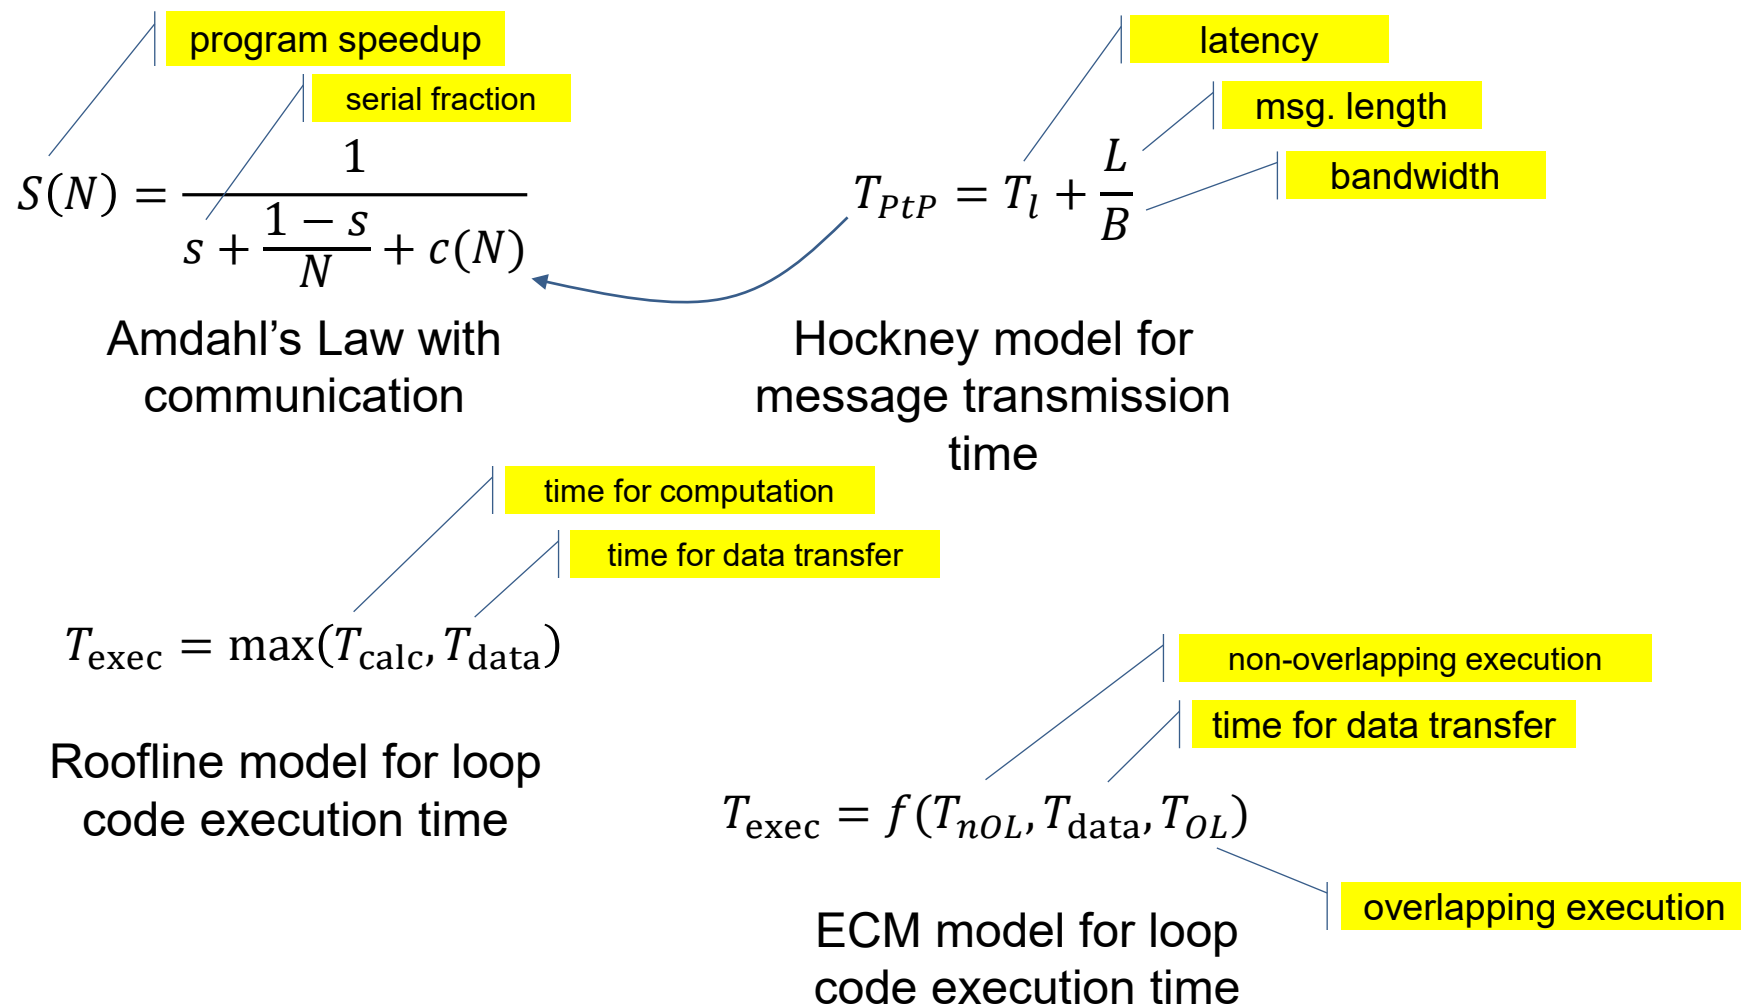
- Curve-fitting analytic model

} black
box

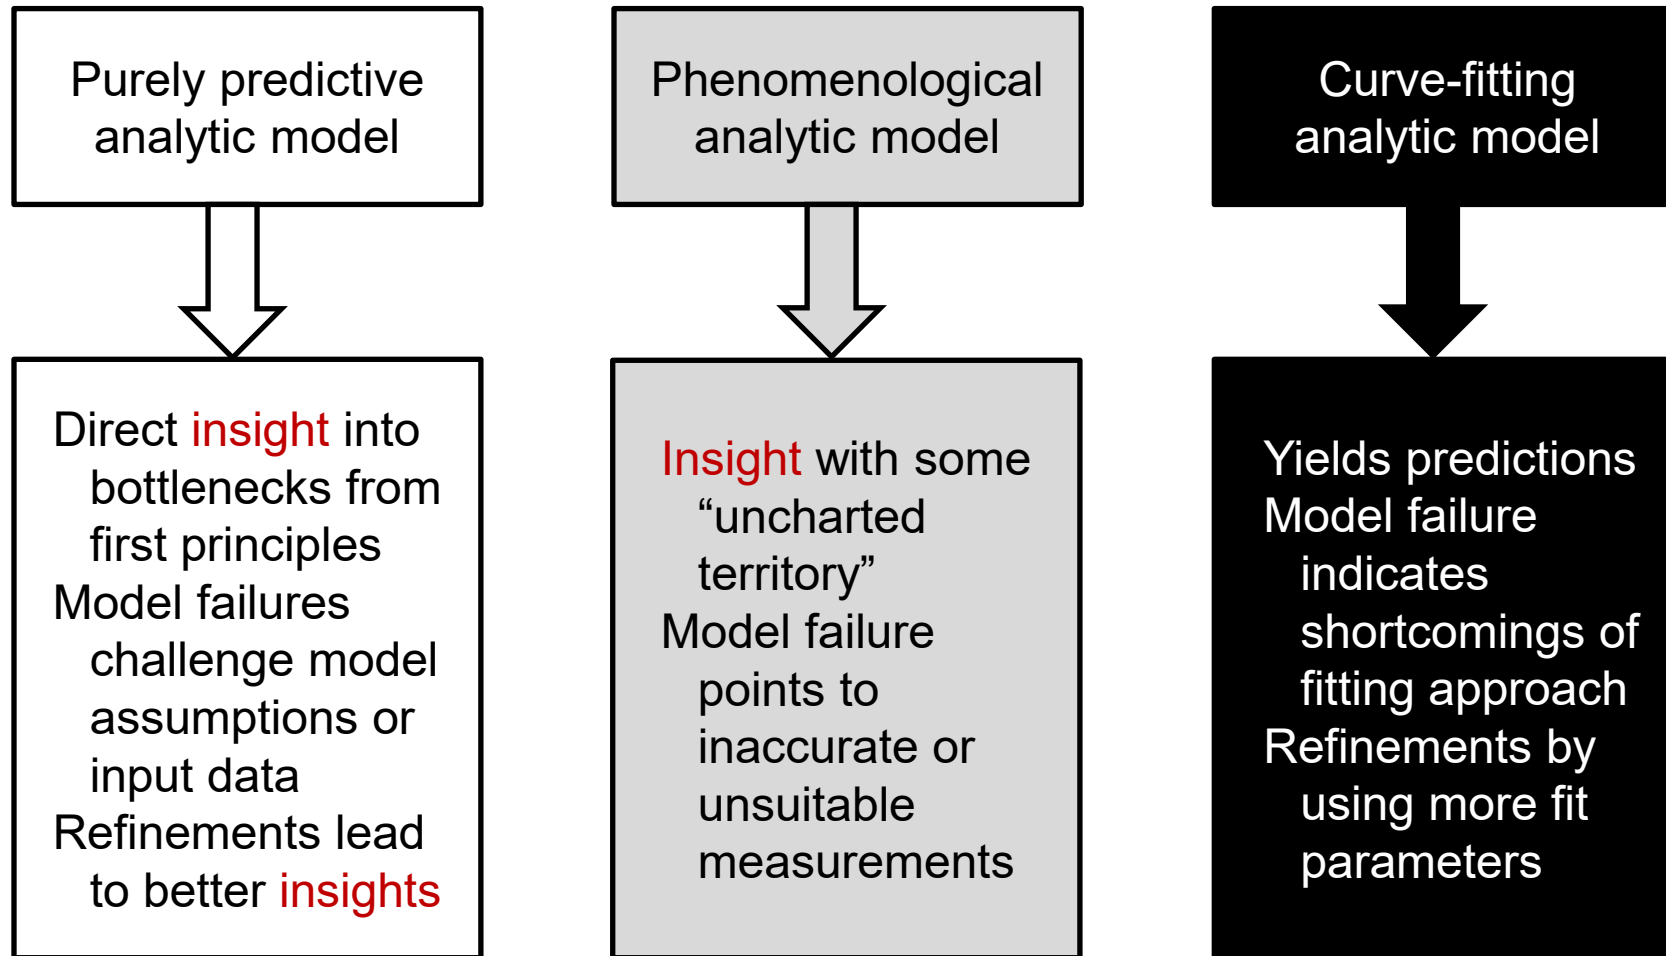
An example from physics



Examples for white-/gray-box models in computing



Models and insights



Resource-based performance models

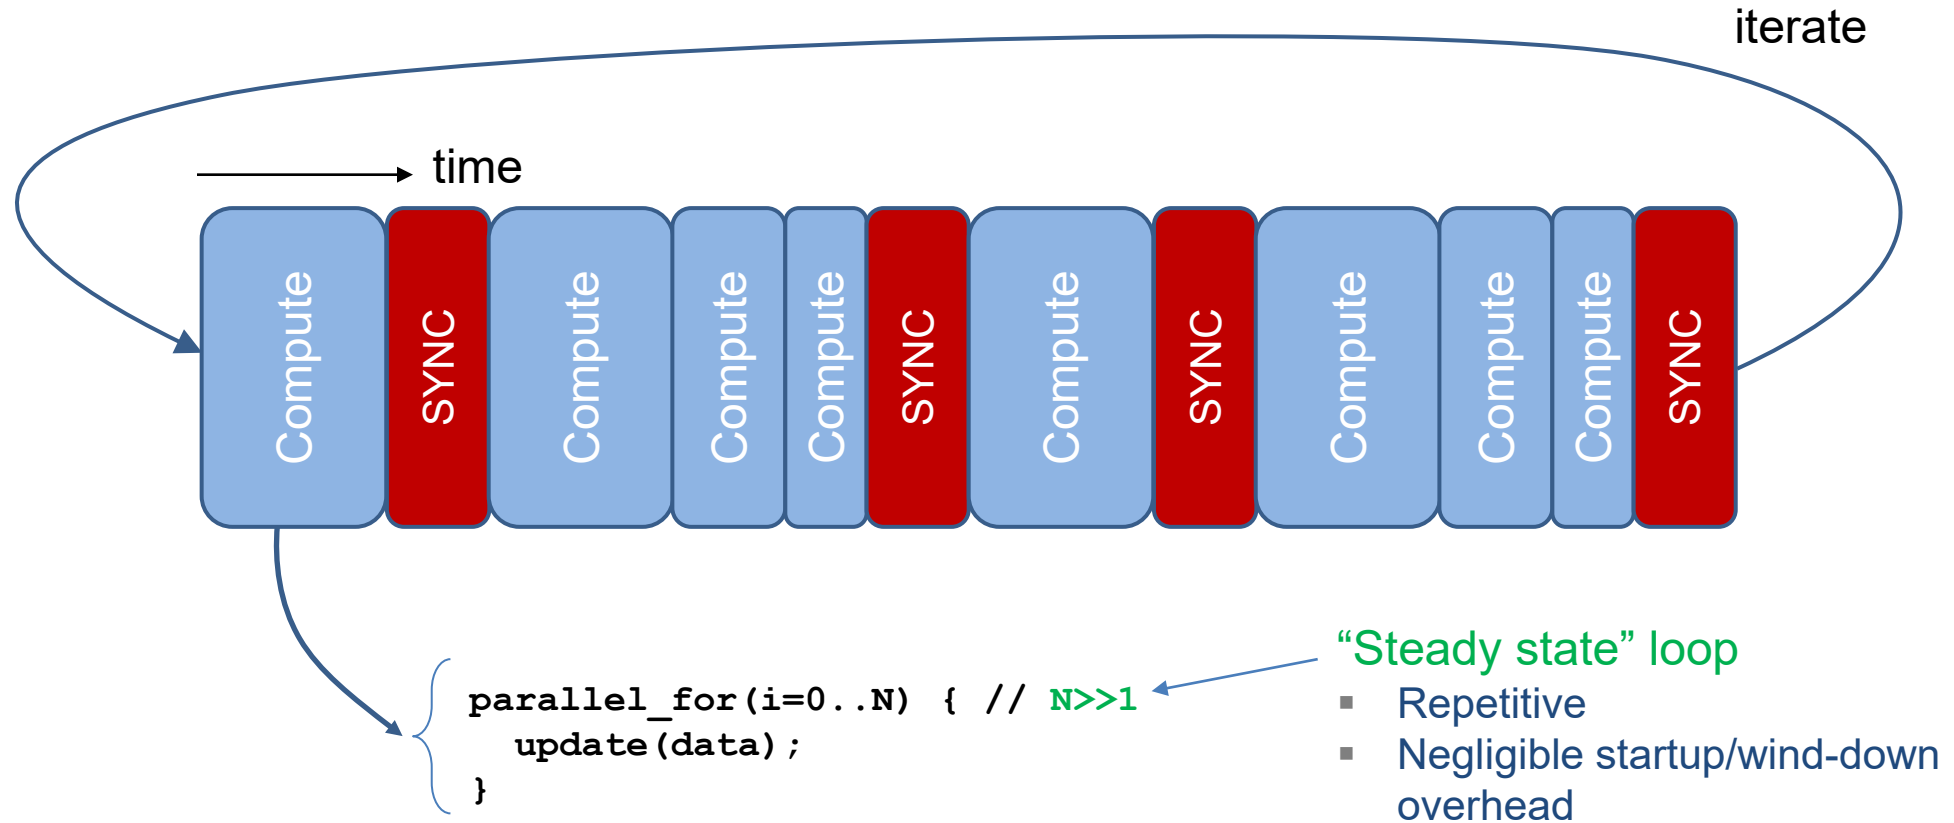


The questions we ask

How much of **\$RESOURCE** does **\$STUFF**
need on **\$HARDWARE**, and why?

→ Analytic, resource-based,
first-principles models

Structure of a solver code




Runtime model: $T = f(\$STUFF, \$HARDWARE)$

Mechanistic vs. resource-based modeling

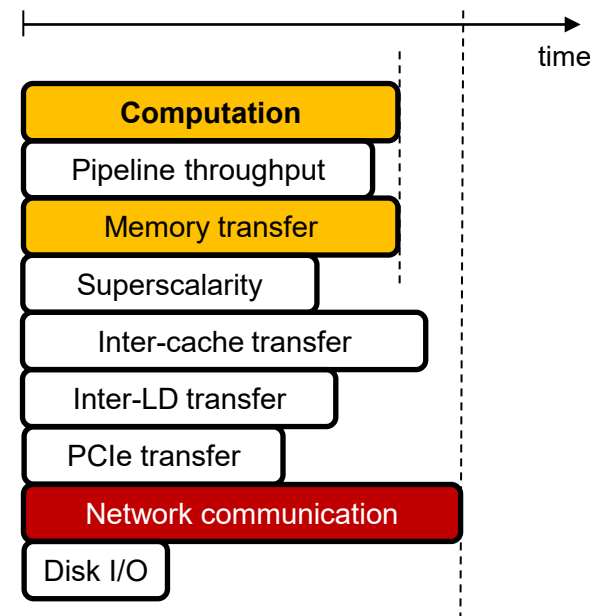
Mechanistic

- Cycle-by-cycle
- Latency
- Simulators
- “It’s complicated”



Resource based

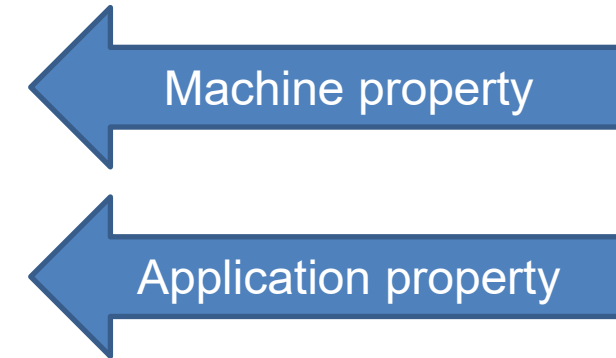
- Resource utilization
- Data flow
- (Non-)overlapping components
- Simplify to make manageable
- If it doesn’t work, refine and iterate




A general view on resource bottlenecks

What is the **maximum performance** when limited by a **bottleneck**?

- Resource bottleneck i delivers resources at maximum rate R_i^{max}
- W_i = needed amount of resources



- **Minimum runtime:** $T_i = \frac{W_i}{R_i^{max}} + \lambda_i$



Bottlenecks

Bottlenecked resources?

- Arithmetic (computational) throughput
- Overall instruction throughput
- Register-L1 data access
- Cache data transfer
- Memory data transfer
- Inter-ccNUMA domain data transfer
- Network data transfer
- ...

Multiple bottlenecks

- Multiple bottlenecks → multiple min. runtimes:

$$T_{\text{expect}} = f(T_1, \dots, T_n)$$



- Overall performance:

$$P_{\text{expect}} = \frac{W}{T_{\text{expect}}}$$

A bottleneck model of computing


Example: two bottlenecks

```
#pragma omp parallel for
for(i=0; i<107; ++i)
    a[i] = a[i] + s * c[i];
```

$$W_{flops} = 2 \times 10^7 \text{ flops}$$

$$W_{BW} = 3 \times 8 \times 10^7 \text{ bytes}$$

$$T_{flops} = \frac{2 \times 10^7 \text{ flops}}{768 \frac{\text{Gflops}}{\text{s}}} = 26.0 \mu\text{s}$$



$$T_{BW} = \frac{2.4 \times 10^8 \text{ bytes}}{210 \frac{\text{Gbyte}}{\text{s}}} = 1.14 \text{ ms}$$

Bottleneck models

How do we reconcile the multiple bottlenecks?
I.e., what is the functional form of $f(T_1, \dots, T_n)$?

- pessimistic model (no overlap): $f(T_1, \dots, T_n) = \sum_i T_i$
- optimistic model (full overlap): $f(T_1, \dots, T_n) = \max(T_1, \dots, T_n)$

Roofline model
(Hockney et al., 1989
Williams et al., 2008)

Roofline for our example: $T_{\min} = \max(T_{flops}, T_{BW}) = 1.14 \text{ ms}$


Maximum performance (“light speed”): $P_{\text{expect}} = \frac{2 \times 10^7}{1.14 \times 10^{-3}} \frac{\text{flops}}{\text{s}} = 17.5 \text{ Gflop/s}$

The Roofline Model



Deriving Roofline from the generic bottleneck model

- Two bottlenecks only
 - Peak computational rate R_{flops}^{max}
 - Peak memory bandwidth R_{BW}^{max}
- Full overlap assumption!
- Runtime prediction:

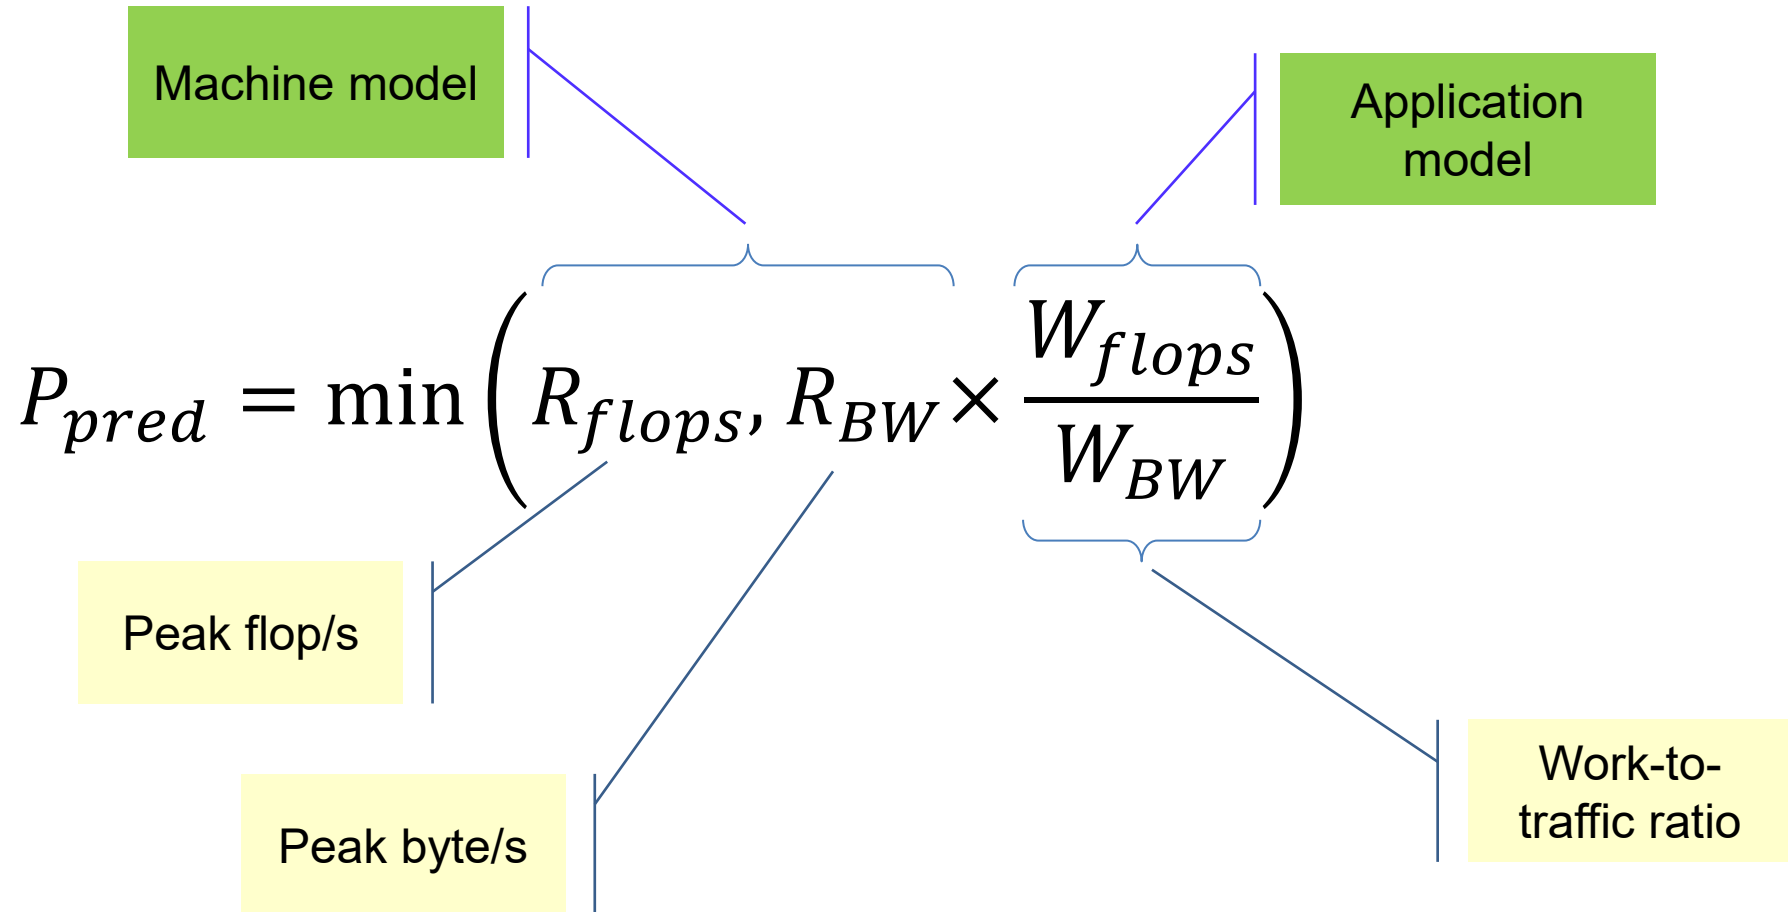


$$T_{pred}(T_{flops}, T_{BW}) = \max(T_{flops}, T_{BW}) = \max\left(\frac{W_{flops}}{R_{flops}^{max}}, \frac{W_{BW}}{R_{BW}^{max}}\right)$$

- Performance prediction in flop/s:


$$P_{pred} = \frac{W_{flops}}{T_{pred}} = \frac{W_{flops}}{\max\left(\frac{W_{flops}}{R_{flops}^{max}}, \frac{W_{BW}}{R_{BW}^{max}}\right)} = \min\left(R_{flops}^{max}, R_{BW}^{max} \times \frac{W_{flops}}{W_{BW}}\right)$$

The two-bottleneck model



The conventional form of the two-bottleneck model

Simplistic view of the hardware:



Simplistic view of the software:

```
! may be multiple levels
do i = 1,<sufficient>
  <complicated stuff doing
    N flops causing
    V bytes of data transfer>
enddo
```

Computational intensity $I = \frac{N}{V}$
→ Unit: flop/byte

Naïve Roofline Model

How fast can tasks be processed? P [flop/s]


The bottleneck is either

- The execution of work: P_{peak} [flop/s]
- The data path: $I \cdot b_S$ [flop/byte x byte/s]

$$P = \min(P_{\text{peak}}, I \cdot b_S)$$

This is the “Naïve Roofline Model”

- High intensity: P limited by execution
- Low intensity: P limited by data transfer
- “Knee” at $P_{\text{peak}} = I \cdot b_S$:
Best use of resources
- Roofline is an “optimistic” model
(think “light speed”)



The Roofline Model in computing – Basics

Apply the naive Roofline model in practice


- Machine parameter #1: Peak performance: $P_{peak} \left[\frac{F}{s} \right]$
 - Machine parameter #2: Memory bandwidth: $b_S \left[\frac{B}{s} \right]$
 - Code characteristic: Computational intensity: $I \left[\frac{F}{B} \right]$
- Machine model
- Application model

Machine properties:

$$P_{peak} = 4 \frac{GF}{s}$$

$$b_S = 10 \frac{GB}{s}$$

Application property: I



Prerequisites for the Roofline Model


- Data transfer and core execution overlap perfectly!
 - Either the limit is core execution or it is data transfer
- Slowest limiting factor “wins”; all others are assumed to have no impact
 - If two bottlenecks are “close,” no interaction is assumed
- Data access latency is ignored, i.e. perfect streaming mode
 - Achievable bandwidth is the limit
- Chip must be able to saturate the bandwidth bottleneck(s)
 - Always model the full chip




Roofline for architecture and code comparison

With Roofline, we can

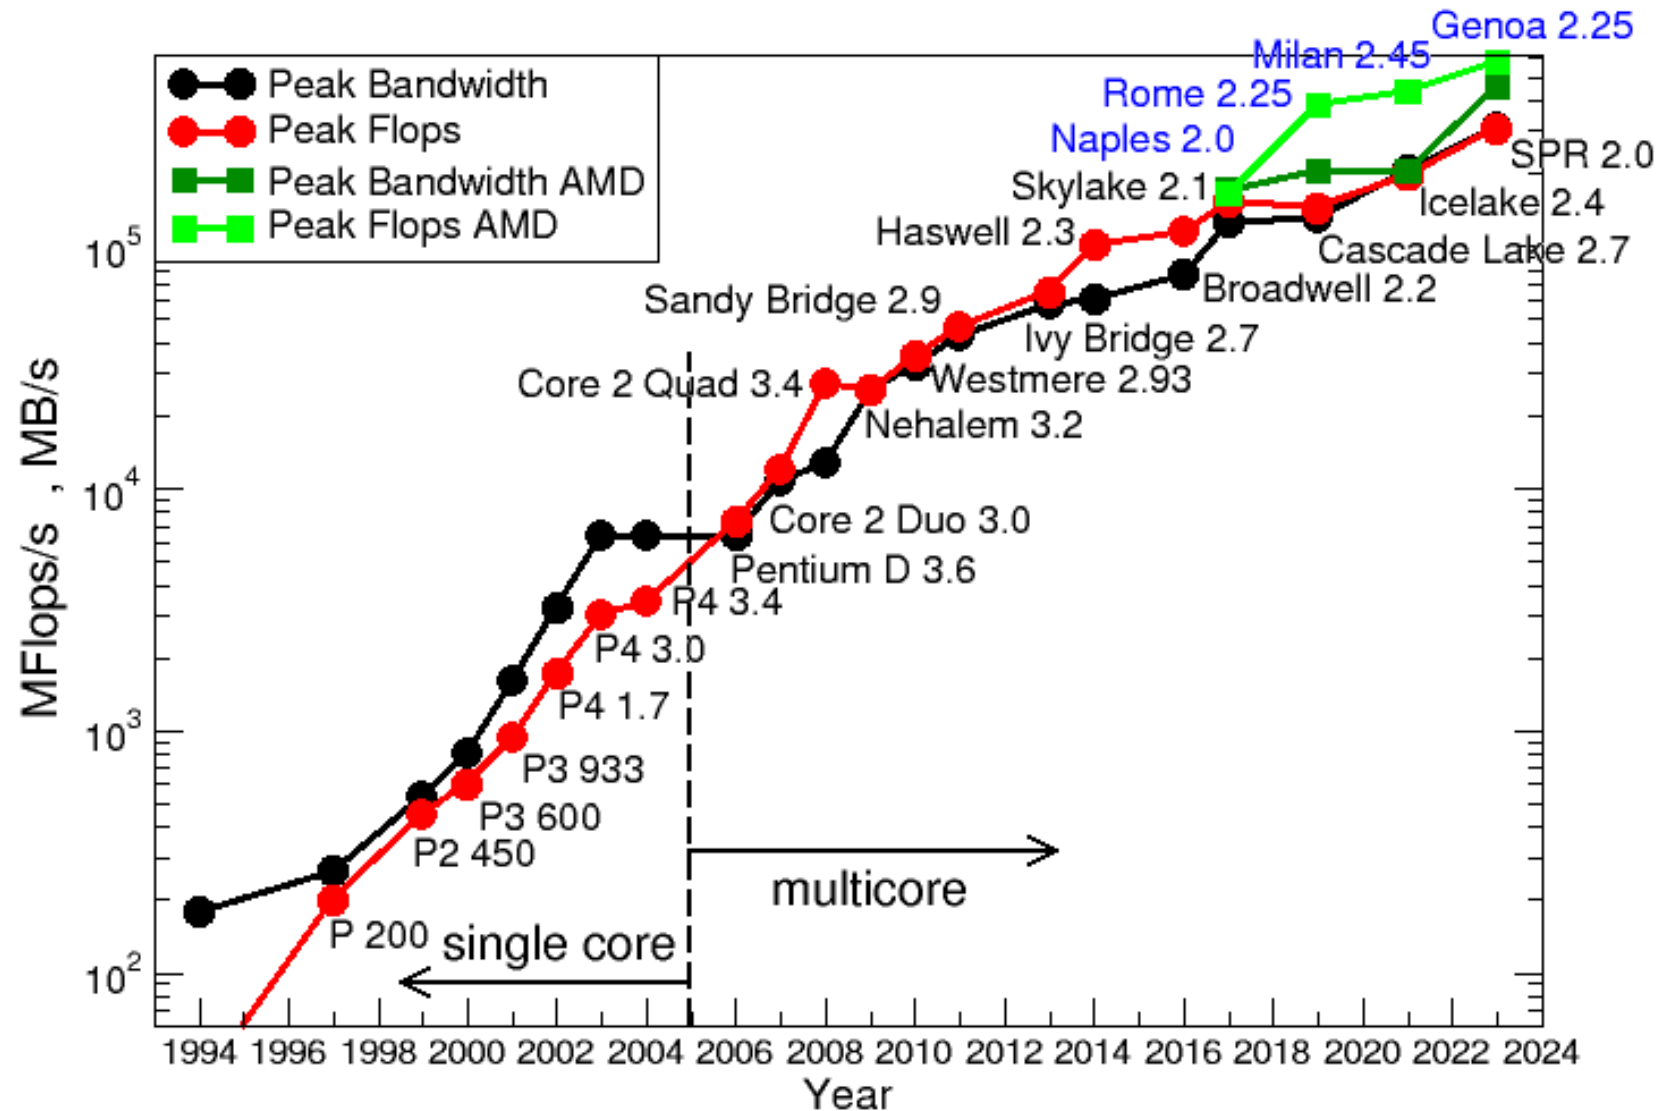
- Compare capabilities of different machines
- Compare performance expectations for different loops
- Roofline always provides upper bound – but is it realistic?
 - Simple case: Loop kernel has loop-carried dependencies → cannot achieve peak
 - Other bandwidth bottlenecks may apply



The “DRAM Gap” (Intel and AMD CPUs)



A different view of the DRAM gap: no SIMD, no FMA



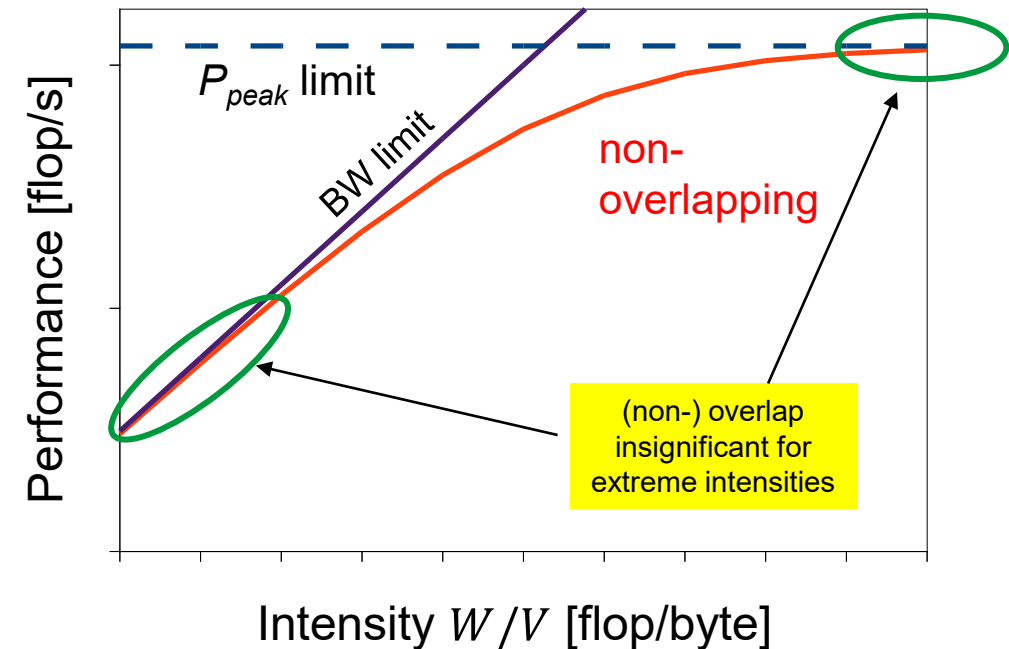
Interlude: a question

What would happen if we switched from the (optimistic) full overlap to the (pessimistic) non-overlap assumption in the two-bottleneck model?

$$T_{pred}(T_{flops}, T_{BW}) = \max(T_{flops}, T_{BW})$$



$$T_{pred}(T_{flops}, T_{BW}) = T_{flops} + T_{BW}$$



A refined Roofline Model

1. P_{\max} = Applicable peak performance of a loop, assuming that data comes from the level 1 cache (this is **not necessarily** P_{peak})
→ e.g., $P_{\max} = 176$ GFlop/s
2. b_S = Applicable (saturated) peak bandwidth of the slowest data path utilized
→ e.g., $b_S = 56$ GByte/s
3. I = Computational intensity (“work” per byte transferred) over the slowest data path utilized (code balance $B_C = I^{-1}$)
→ e.g., $I = 0.167$ Flop/Byte → $B_C = 6$ Byte/Flop

“Flop” is not the only
useful unit of work!

Performance limit:

$$P = \min(P_{\max}, I \cdot b_S) = \min\left(P_{\max}, \frac{b_S}{B_C}\right)$$

[byte/s]

[byte/Flop]

A “simple” example: The sum reduction



Machine and code characteristics

- Clock speed: 2.2 GHz
- AVX-capable (256-bit SIMD)
- 1 ADD, 1 MUL instruction per cy
- 8 cores
- Single-precision peak:
$$P_{peak} = 8 \times 2 \times 8 \times 2.2 \text{ Gflop/s}$$
$$= 281.6 \text{ Gflop/s}$$
- ADD pipeline latency: 3 cy
- Memory bandwidth $b_S = 40 \text{ Gbyte/s}$

```
float sum, a[N];  
//...  
#pragma omp parallel for \  
        reduction(+:sum)  
for (int i=0; i<N; i++){  
    sum += a[i];  
}
```

Arithmetic intensity:


$$I = \frac{1 \text{ flop}}{4 \text{ byte}} = 0.25 \text{ F/B}$$

Code balance:

$$B_c = 4 \text{ B/F}$$

Applicable peak performance

```
for (int i=0; i<N; i++) {  
    sum += a[i];  
}
```



How fast can this loop possibly run with data in the L1 cache? (P_{max})

- **Loop-carried dependency** on summation variable
- Execution **stalls** at every ADD until previous ADD is complete

→ No pipelining?

→ No SIMD?

Applicable peak for the sum reduction (I)

Plain scalar code, no SIMD


```
for (int i=0; i<N; i++){  
    sum += a[i];  
}
```

LOAD r1.0 ← 0
i ← 1
loop:
 LOAD r2.0 ← a(i)
 ADD r1.0 ← r1.0 + r2.0
 ++i →? loop
result ← r1.0

SIMD lane

SIMD lanes

ADD pipes utilization:



→ 1/24 of ADD peak

Applicable peak for the sum reduction (II)

Scalar code, 3-way “modulo variable expansion”

```
LOAD r1.0 ← 0
LOAD r2.0 ← 0
LOAD r3.0 ← 0
i ← 1
```

loop:

```
LOAD r4.0 ← a(i)
LOAD r5.0 ← a(i+1)
LOAD r6.0 ← a(i+2)
```

```
ADD r1.0 ← r1.0 + r4.0 # scalar ADD
ADD r2.0 ← r2.0 + r5.0 # scalar ADD
ADD r3.0 ← r3.0 + r6.0 # scalar ADD
```

```
i+=3 →? loop
result ← r1.0+r2.0+r3.0
```

```
for (int i=0; i<N; i+=3){
    s1 += a[i+0];
    s2 += a[i+1];
    s3 += a[i+2];
}
sum = sum + s1+s2+s3;
```

s1	s2	s3

→ 1/8 of ADD peak

Applicable peak for the sum reduction (III)

SIMD vectorization (8-way MVE) x
pipelining (3-way MVE)

```
LOAD [r1.0,...,r1.7] ← [0,...,0]
LOAD [r2.0,...,r2.7] ← [0,...,0]
LOAD [r3.0,...,r3.7] ← [0,...,0]
i ← 1
```

loop:


```
LOAD [r4.0,...,r4.7] ← [a(i),...,a(i+7)] # SIMD LOAD
LOAD [r5.0,...,r5.7] ← [a(i+8),...,a(i+15)] # SIMD
LOAD [r6.0,...,r6.7] ← [a(i+16),...,a(i+23)] # SIMD
```

```
ADD r1 ← r1 + r4 # SIMD ADD
ADD r2 ← r2 + r5 # SIMD ADD
ADD r3 ← r3 + r6 # SIMD ADD
```

i+=24 →? loop

result ← r1.0+r1.1+...+r3.6+r3.7

```
for (int i=0; i<N; i+=24){
    s10 += a[i+0]; s20 += a[i+8]; s30 += a[i+16];
    s11 += a[i+1]; s21 += a[i+9]; s31 += a[i+17];
    s12 += a[i+2]; s22 += a[i+10]; s32 += a[i+18];
    s13 += a[i+3]; s23 += a[i+11]; s33 += a[i+19];
    s14 += a[i+4]; s24 += a[i+12]; s34 += a[i+20];
    s15 += a[i+5]; s25 += a[i+13]; s35 += a[i+21];
    s16 += a[i+6]; s26 += a[i+14]; s36 += a[i+22];
    s17 += a[i+7]; s27 += a[i+15]; s37 += a[i+23];
}
sum = sum + s10+s11+...+s37;
```



Sum reduction

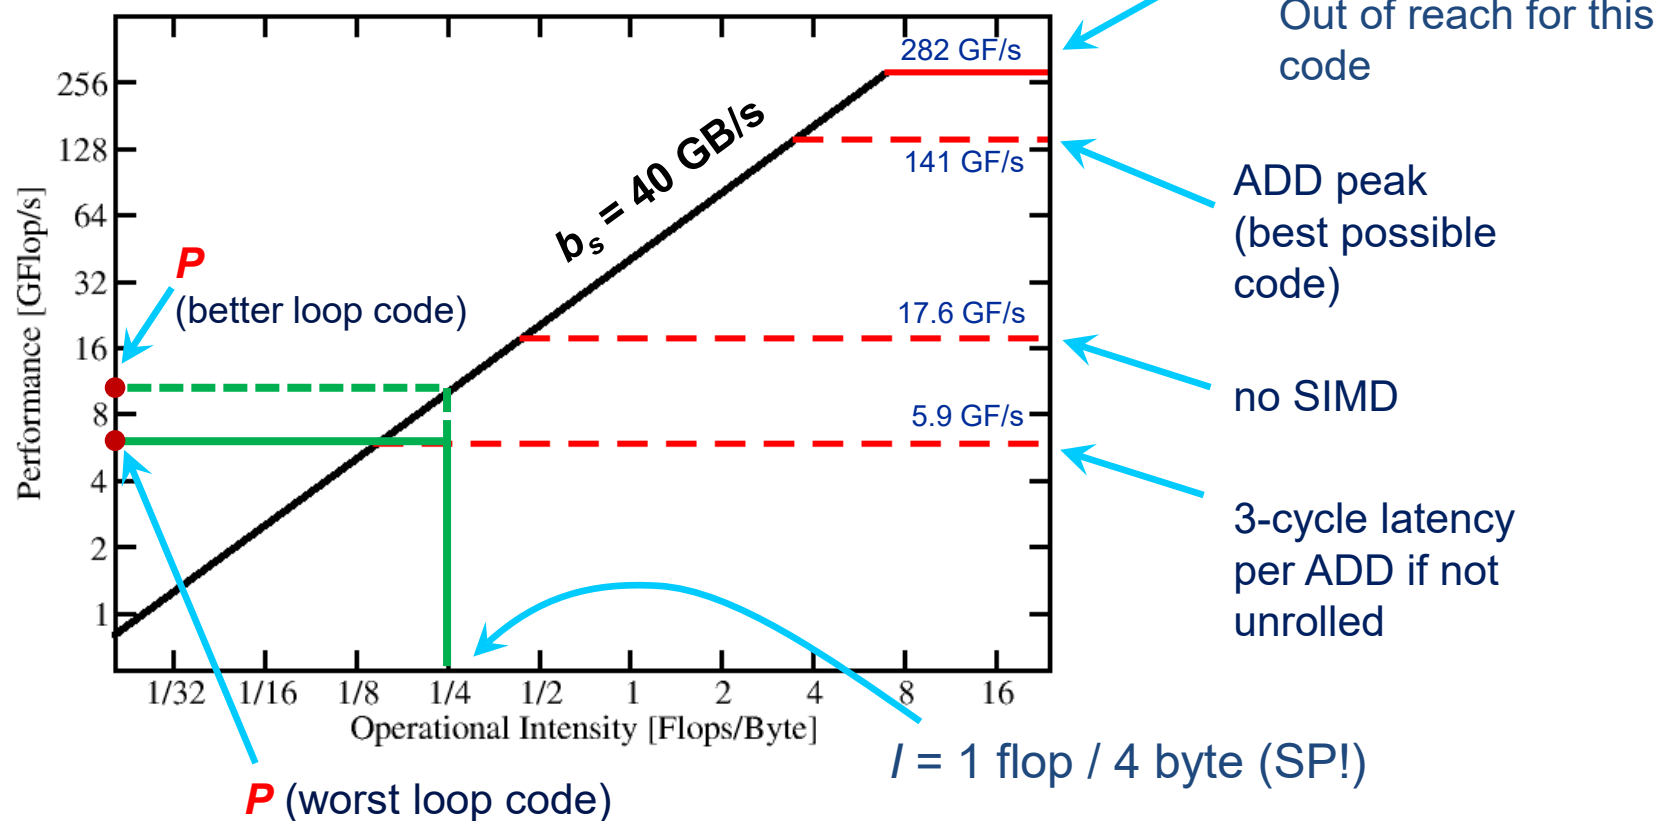
Questions

- When can this performance actually be achieved?
 - No **data transfer** bottlenecks
 - No other **in-core** bottlenecks
 - Need to execute (3 LOADs + 3 ADDs + 1 increment + 1 compare + 1 branch) in 3 cycles
- What does the **compiler** do?
 - If allowed and capable, the compiler will do this automatically
- Is the compiler **allowed** to do this at all?
 - Not according to language standards
 - High optimization levels can violate language standards
- What about the “accuracy” of the result?
 - Good question ;-)

Now let's add the bandwidth ceiling

```
float sum, a[N];
//...
#pragma omp parallel for \
    reduction(+:sum)
for (int i=0; i<N; i++){
    sum += a[i];
}
```

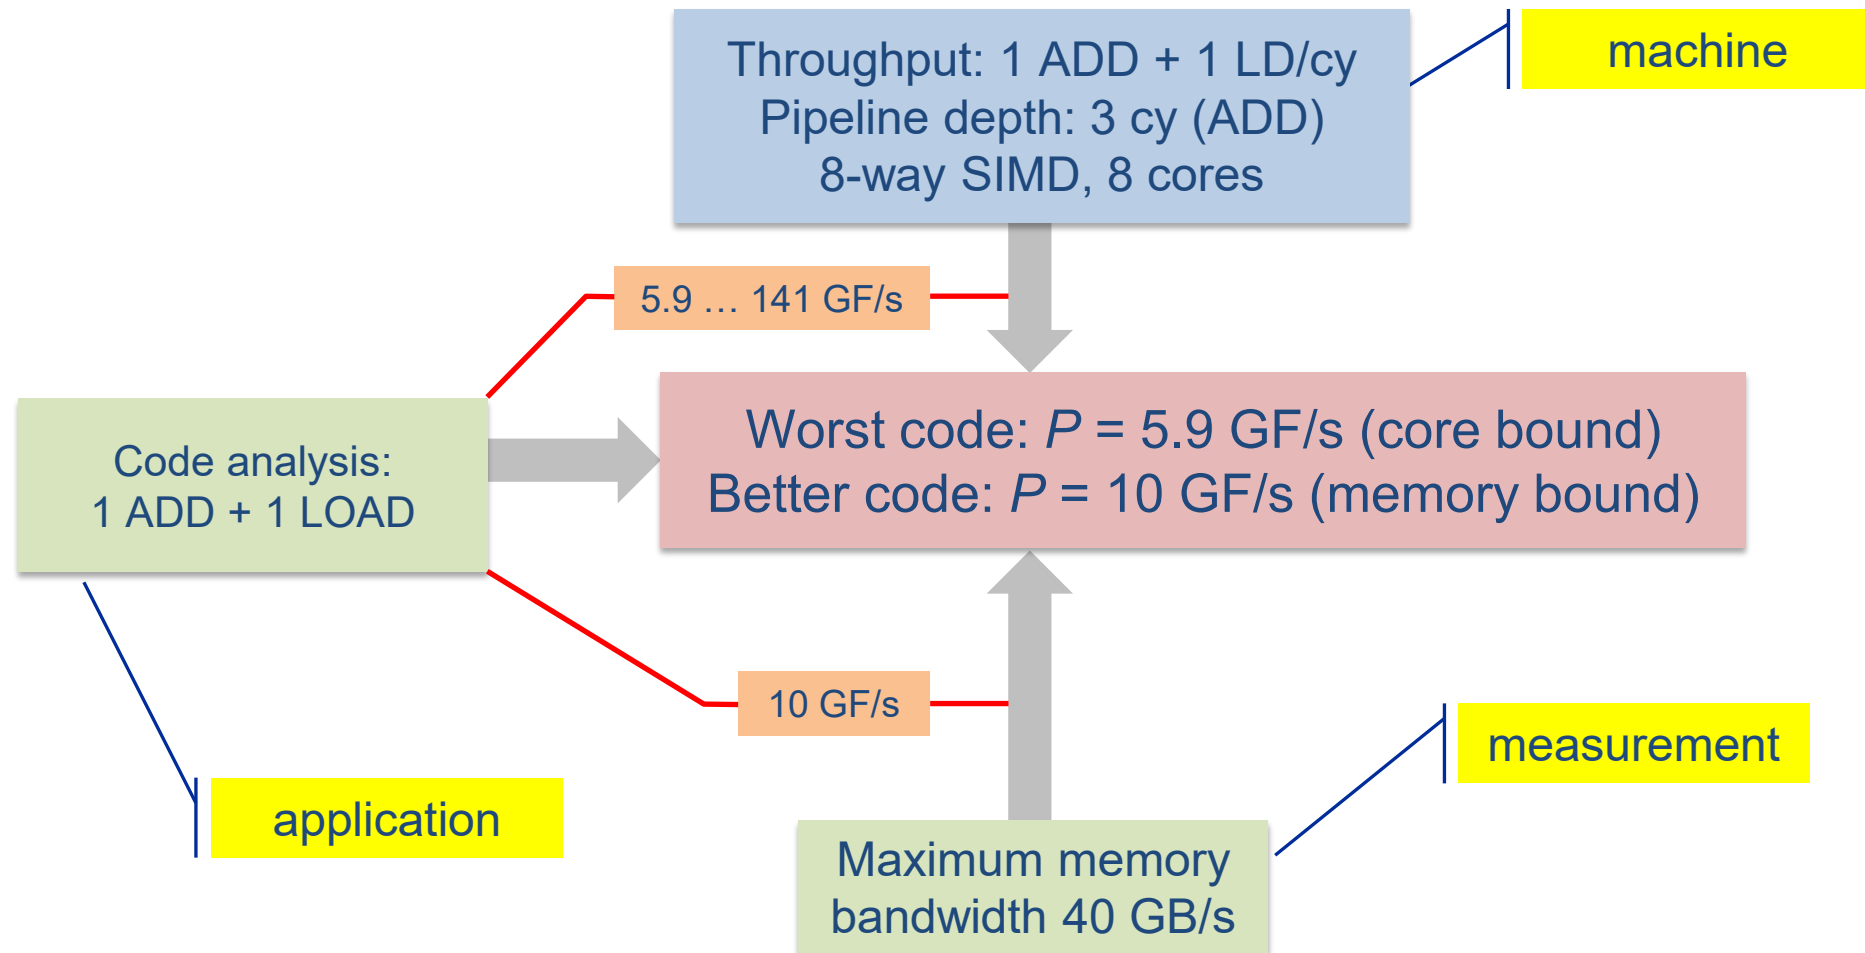
$$P = \min(P_{\max}, I \cdot b_s)$$



Input to the roofline model

... on the example of
in single precision


```
do i=1,N; s=s+a(i); enddo
```



Factors to consider in the Roofline Model


Bandwidth-bound

1. Accurate traffic calculation (write-allocate, strided access, cache reuse,...)
2. Practical \neq theoretical BW limits
3. Saturation effects \rightarrow consider full socket only



Core-bound

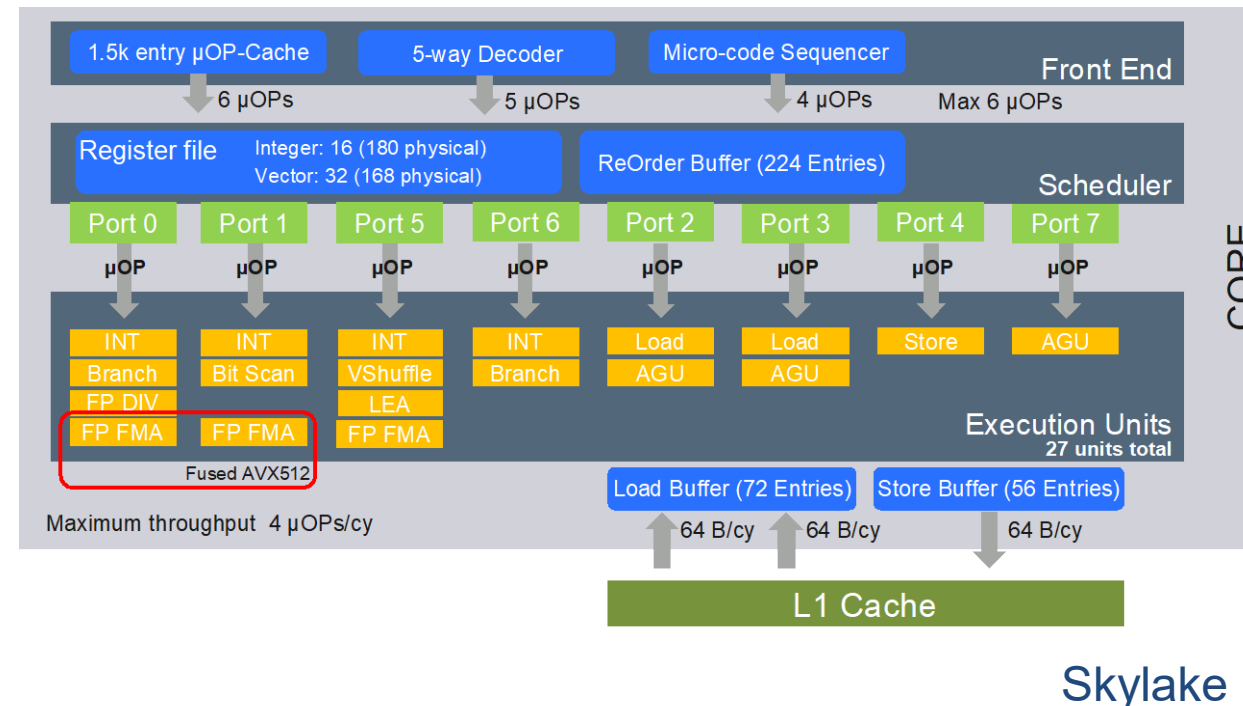
1. Multiple bottlenecks: LD/ST, arithmetic, pipelines, SIMD, execution ports
2. Limit is linear in # of cores



Complexities of in-core execution (P_{\max})

Multiple bottlenecks:

- Decode/retirement throughput
- Port contention (direct or indirect)
- Arithmetic pipeline stalls (dependencies)
- Overall pipeline stalls (branching)
- L1 Dcache bandwidth (LD/ST throughput)
- Scalar vs. SIMD execution
- L1 Icache (LD/ST) bandwidth
- Alignment issues
- ...



Tool for P_{\max} analysis: OSACA

<http://tiny.cc/OSACA>

DOI: [10.1109/PMBS49563.2019.00006](https://doi.org/10.1109/PMBS49563.2019.00006)

DOI: [10.1109/PMBS.2018.8641578](https://doi.org/10.1109/PMBS.2018.8641578)

Hardware features of (some) Intel Xeon processors

Microarchitecture	Ivy Bridge EP	Broadwell EP	Cascade Lake SP	Ice Lake SP
Introduced	09/2013	03/2016	04/2019	06/2021
Cores	≤ 12	≤ 22	≤ 28	≤ 40
LD/ST throughput per cy:				
AVX(2), AVX512	1 LD + ½ ST	2 LD + 1 ST	2 LD + 1 ST	2 LD + 1 ST
SSE/scalar	2 LD 1 LD & 1 ST			
ADD throughput	1 / cy	1 / cy	2 / cy	2 / cy
MUL throughput	1 / cy	2 / cy	2 / cy	2 / cy
FMA throughput	N/A	2 / cy	2 / cy	2 / cy
L1-L2 data bus	32 B/cy	64 B/cy	64 B/cy	64 B/cy
L2-L3 data bus	32 B/cy	32 B/cy	16+16 B/cy	16+16 B/cy
L1/L2 per core	32 KiB / 256 KiB	32 KiB / 256 KiB	32 KiB / 1 MiB	48 KiB / 1.25 MiB
LLC	2.5 MiB/core inclusive	2.5 MiB/core inclusive	1.375 MiB/core exclusive/victim	1.5 MiB/core exclusive/victim
Memory	4ch DDR3	4ch DDR3	6ch DDR4	8ch DDR4
Memory BW (meas.)	~ 48 GB/s	~ 62 GB/s	~ 115 GB/s	~ 160 GB/s

Source: <https://software.intel.com/content/www/us/en/develop/download/intel-64-and-ia-32-architectures-optimization-reference-manual.html>

Code balance: more examples

```
double a[], b[];
for(i=0; i<N; ++i)
    a[i] = a[i] + b[i];
```

$$B_C = 24B / 1F = 24 \text{ B/F}$$
$$I = 0.042 \text{ F/B}$$

```
double a[], b[];
for(i=0; i<N; ++i)
    a[i] = a[i] + s * b[i];
```

$$B_C = 24B / 2F = 12 \text{ B/F}$$
$$I = 0.083 \text{ F/B}$$

```
float s=0, a[];
for(i=0; i<N; ++i)
    s = s + a[i] * a[i];
```

Scalar – can be kept in register

$$B_C = 4B / 2F = 2 \text{ B/F}$$
$$I = 0.5 \text{ F/B}$$

```
float s=0, a[], b[];
for(i=0; i<N; ++i)
    s = s + a[i] * b[i];
```

Scalar – can be kept in register

$$B_C = 8B / 2F = 4 \text{ B/F}$$
$$I = 0.25 \text{ F/B}$$

```
float s=0, a[], b[];
for(i=0; i<N; ++i)
    for(j=0; j<N; ++j)
        b[i][j] = a[i][j]
                + a[i-1][j]
                + a[i+1][j];
```

Scalar – can be kept in register

$$B_C = 16B / 2F \text{ or}$$
$$8B / 2F \text{ or even } ???$$
$$20 \text{ B} / 2F$$

Streaming, perfect spatial locality, no temporal locality
→ simple

And what about this?

```
float s=0, a[N], b[N];
int idx[N];
for(i=0; i<N; ++i)
    s = s + a[i]
        * b[idx[i]];
```

Possible cache reuse → tricky!

We'll get to it!

Is there anything to ease the construction of the model?

Code balance B_C


- Close inspection and hard thinking
- Simplifying assumptions
 - “What is the minimum possible amount of traffic?”
 - “What is the worst case?”
- Tools
 - Kerncraft
<https://github.com/RRZE-HPC/kerncraft>

In-core P_{\max}

- Inspection of assembly code and manual modeling
- Simplifying assumptions
 - “What is the required minimum number of arithmetic/load/store instructions?”
 - $P_{\max} = P_{peak}$
- Tools
 - OSACA
<https://github.com/RRZE-HPC/OSACA>


Tracking code optimizations in the Roofline Model

1. Hit the BW bottleneck by good serial code
(e.g., Python \rightarrow Fortran)
2. Increase intensity to make better use of BW bottleneck
(e.g., spatial loop blocking)
3. Increase intensity and go from memory bound to core bound
(e.g., temporal blocking)
4. Hit the core bottleneck by good serial code
(e.g., `-fno-alias`, SIMD intrinsics)



Shortcomings of the roofline model

- Saturation effects in multicore chips are not explained
 - Reason: “saturation assumption”
 - Cache line transfers and core execution do sometimes not overlap perfectly
 - It is not sufficient to measure single-core STREAM to make it work
 - Only increased “pressure” on the memory interface can saturate the bus
→ need more cores!
- In-cache performance is not accurately predicted
- The ECM performance model gives more insight:



G. Hager, J. Treibig, J. Habich, and G. Wellein: Exploring performance and power properties of modern multicore chips via simple machine models. Concurrency and Computation: Practice and Experience (2013).


DOI: [10.1002/cpe.3180](https://doi.org/10.1002/cpe.3180) Preprint: [arXiv:1208.2908](https://arxiv.org/abs/1208.2908)

Multi-ceiling Roofline model: graphical representation

Multiple ceilings may apply

- Different **bandwidths / data paths**
→ different inclined ceilings
→ possibly different I (and B_c) for the same kernel
- Different P_{\max}
→ different flat ceilings

In fact, P_{\max} should always come from **code analysis**; generic ceilings are usually impossible to attain



What about multiple loops (i.e., solvers)?

Performance-based formulation is inadequate → go back to time

Solver: s components $j = 1 \dots s$, t_j = model time for component j

$$t_{solver} = \sum_{j=1}^s t_j = \sum_{j=1}^s f(T_{1,j}, \dots, T_{N,j})$$

“Roofline”:

$$t_{solver} = \sum_{j=1}^s \max(T_{flops,j}, T_{BW,j})$$

$$P_{solver} = \frac{\text{Whatever—you—consider—a—good—metric—for—work}}{t_{solver}}$$

What if only the bandwidth bottleneck applies?

All loops memory bound \rightarrow runtime entirely determined by data transfer

$$t_{solver} = \sum_{j=1}^s \max\left(\frac{W_{flops,j}}{R_{flops,j}}, \frac{W_{BW,j}}{R_{BW,j}}\right) = \sum_{j=1}^s \frac{W_{BW,j}}{R_{BW,j}} = \frac{1}{R_{BW}} \sum_{j=1}^s W_{BW,j} = \frac{V}{b_S}$$

Overall data
volume V

Overall “work”


$$P_{solver} = \frac{W}{t_{solver}} = \frac{W}{V} \times b_S = I \times b_S$$

Diagnostic / phenomenological Roofline modeling





Diagnostic modeling

- What if we cannot predict the intensity/balance?
 - Code very complicated
 - Code not available
 - Parameters unknown
 - Doubts about correctness of analysis
- Measure data volume V_{meas} (and work N_{meas})
 - Hardware performance counters
 - Tools: likwid-perfctr, PAPI, Intel Vtune,...
- Insights + benefits
 - Compare analytic model and measurement → validate model
 - Can be applied (semi-)automatically
 - Useful in performance monitoring of user jobs on clusters



Roofline and performance monitoring of clusters

Two cluster jobs...




Which of them is
“good” and which is
“bad”?



Diagnostic modeling of a complex code (3 kernels)

Multiple bandwidth bottlenecks

→ need I for each one ($I_{mem}, I_{L3}, I_{L2}, \dots$)



Kernel 1 ●

- Performance close to memory BW ceiling but far away from others
→ indicates **memory bound**

Kernel 2 ▲

- Performance not near any BW ceiling
- Performance close to scalar peak ceiling
→ indicates **scalar core-bound peak code**

Kernel 3 ■

- Performance not anywhere near any ceiling
→ There must be an (as yet) **unknown in-core performance limit** P_{max}

Roofline conclusion

- Roofline = simple, first-principle, resource-based model for upper performance limit of data-streaming loops
 - Machine model (P_{max}, b_S, \dots) + application model (I_{mem}, I_{L3}, \dots)
 - Conditions apply, extensions exist
- Two modes of operation; per kernel:
 - Predictive: Calculate I_j , calculate upper limit, validate model, optimize, iterate
 - Diagnostic: Measure I_j and P , compare with ceilings
- Challenge of predictive modeling: Getting P_{max} and I_* right

Performance Analysis with Hardware Metrics

... on the example of likwid-perfctr



Tools for Node-level Performance Engineering

- **Node Information**

*/proc/cpuinfo, numactl, hwloc, **likwid-topology**, likwid-powermeter*

- **Affinity control** and data placement

*OpenMP and MPI runtime environments, hwloc, numactl, **likwid-pin***

- **Runtime Profiling**

Compilers, gprof, perf, HPC Toolkit, Intel Amplifier, ...

- **Performance Analysis**

*Intel VTune, **likwid-perfctr**, PAPI-based tools, HPC Toolkit, Linux perf*

- **Microbenchmarking**

*STREAM, **likwid-bench**, Imbench, uarch-bench*

LIKWID performance tools

LIKWID tool suite:

Like
I
Knew
What
I'm
Doing




<https://youtu.be/6uF11HPq-88>

Open source tool collection
(developed at NHR@FAU):



<https://github.com/RRZE-HPC/likwid>



J. Treibig, G. Hager, G. Wellein: *LIKWID: A lightweight performance-oriented tool suite for x86 multicore environments*. PSTI2010, Sep 13-16, 2010, San Diego, CA. DOI: [10.1109/ICPPW.2010.38](https://doi.org/10.1109/ICPPW.2010.38)

LIKWID Tool Suite

- Command line tools for Linux:

- easy to install

- works with standard Linux kernel

- simple and clear to use

- supports most X86 CPUs

- (also ARMv8, POWER9 and Nvidia GPUs)



- Current tools:

- likwid-topology** – Print thread and cache topology

- likwid-pin** – Pin threaded application without touching code

- likwid-mpirun** – Start MPI and MPI/OpenMP hybrid programs

- likwid-perfctr** – Measure performance counters

- likwid-bench** – Microbenchmarking tool and environment

... some more

Probing performance behavior


- How do we find out about the performance properties and resource usage of a running code?
- How do we validate analytic (white- or gray-box) performance models?
- How do we get the data for diagnostic RL modeling?

- Many tools exist, but a coarse overview is often sufficient: `likwid-perfctr`
- Simple end-to-end measurement of hardware performance metrics

Operating modes:

- Wrapper
- Stethoscope
- Timeline
- Marker API

Preconfigured and extensible
metric groups, list with
`likwid-perfctr -a`



BRANCH: Branch prediction miss rate/ratio
CACHE: Data cache miss rate/ratio
CLOCK: Clock frequency of cores
DATA: Load to store ratio
FLOPS_DP: Double Precision MFlops/s
FLOPS_SP: Single Precision MFlops/s
L2: L2 cache bandwidth in MBytes/s
L2CACHE: L2 cache miss rate/ratio
L3: L3 cache bandwidth in MBytes/s
L3CACHE: L3 cache miss rate/ratio
MEM: Main memory bandwidth in MBytes/s
TLB: TLB miss rate/ratio
ENERGY: Power and energy consumption

likwid-perfctr wrapper mode

```
$ likwid-perfctr -g L2 -C S1:0-3 ./a.out
```

```
-----  
CPU name: Intel(R) Xeon(R) Platinum 8360Y CPU @ 2.40GHz[...]  
-----
```

```
<<<< PROGRAM OUTPUT >>>>
```

```
Group 1: L2
```

Event	Counter	HWThread 36	HWThread 37	HWThread 38	HWThread 39
INSTR_RETIRED_ANY	FIXC0	1409713380	1393263859	1394342491	1388917034
CPU_CLK_UNHALTED_CORE	FIXC1	2095261718	2088036330	2075539220	2058287996
CPU_CLK_UNHALTED_REF	FIXC2	2103679392	2121235200	2100479808	2075658144
TOPDOWN_SLOTS	FIXC3	10476308590	10440181650	10377696100	10291439980
L1D_REPLACEMENT	PMC0	142720376	142481840	142482162	142434419
L2_TRANS_L1D_WB	PMC1	54986306	54864382	54868339	54815549
TCACHE_64B_IPTAG_MISS	PMC2	381869	2094	7399	7718

Always
measured for
Intel CPUs

Configured metrics
(this group)

```
[... statistics output omitted ...]
```

Metric	HWThread 36	HWThread 37	HWThread 38	HWThread 39
Runtime (RDTSC) [s]	1.0092	1.0092	1.0092	1.0092
Runtime unhaltd [s]	0.8751	0.8721	0.8669	0.8597
Clock [MHz]	2384.7406	2356.8484	2365.8917	2374.2844
CPI	1.4863	1.4987	1.4885	1.4819
L2D load bandwidth [MBytes/s]	9050.5857	9035.4589	9035.4794	9032.4518
L2D load data volume [GBytes]	9.1341	9.1188	9.1189	9.1158
L2D evict bandwidth [MBytes/s]	3486.9462	3479.2144	3479.4653	3476.1177
L2D evict data volume [GBytes]	3.5191	3.5113	3.5116	3.5082
L2 bandwidth [MBytes/s]	12561.7480	12514.8061	12515.4139	12509.0589
L2 data volume [GBytes]	12.6777	12.6303	12.6309	12.6245

Derived
metrics

likwid-perfctr stethoscope mode

- likwid-perfctr counts events on cores; it has no notion of what kind of code is running (if any)



This allows you to “listen” to what is currently happening, **without any overhead:**

```
$ likwid-perfctr -c N:0-11 -g FLOPS_DP -S 10s
```

- Can be used as cluster/server monitoring tool
- Frequent use: monitor a long-running parallel application from outside

Cluster monitoring with likwid-perfctr

Two jobs on the NHR@FAU “Fritz” cluster



<https://github.com/ClusterCockpit>

likwid-perfctr with MarkerAPI

- The MarkerAPI can restrict measurements to **code regions**
- The API only reads counters.
The configuration of the counters is still done by **likwid-perfctr**
- Multiple named regions allowed, accumulation over multiple calls
- Inclusive and overlapping regions allowed
- **Caveat:** Marker API can cause overhead; do not call too frequently!

```
#include <likwid-marker.h>

LIKWID_MARKER_INIT; // must be called from serial region
. . .
LIKWID_MARKER_START("Compute");
. . .
LIKWID_MARKER_STOP("Compute");
. . .
LIKWID_MARKER_START("Postprocess");
. . .
LIKWID_MARKER_STOP("Postprocess");
. . .
LIKWID_MARKER_CLOSE; // must be called from serial region
```

likwid-perfctr with MarkerAPI: OpenMP code (C)

```
#include <likwid-marker.h>

int main(...) {
    LIKWID_MARKER_INIT;
    #pragma omp parallel
    {
        LIKWID_MARKER_REGISTER("MatrixAssembly");
    }
    ...
    #pragma omp parallel
    {
        LIKWID_MARKER_START("MatrixAssembly");
        #pragma omp for
        for(int i=0; i<N; ++i) { /* Loop */ }
        LIKWID_MARKER_STOP("MatrixAssembly");
    }
    ...
    LIKWID_MARKER_CLOSE;
}
```

Optional: Prepare data structures (reduced overhead on 1st marker call)

Call markers in parallel region if data should be taken on all threads

<https://github.com/RRZE-HPC/likwid/wiki/TutorialMarkerC>

likwid-perfctr with MarkerAPI: OpenMP code (Fortran)

```
program p
  use likwid
  call likwid_markerInit
  !$omp parallel
    call likwid_markerRegisterRegion("MatrixAssembly")
  !$omp end parallel
  ...
  !$omp parallel
    call likwid_markerStartRegion("MatrixAssembly")
    !$omp do
      do i=1,N
        ! Loop
      enddo
    !$omp end do
    call likwid_markerStopRegion("MatrixAssembly")
  !$omp end parallel
  ...
  call likwid_markerClose
end program p
```

Optional: Prepare data structures (reduced overhead on 1st marker call)

Call markers in parallel region if data should be taken on all threads

<https://github.com/RRZE-HPC/likwid/wiki/TutorialMarkerF90>

Compiling, linking, and running with marker API

Compile:

```
cc -I /path/to/likwid.h -DLIKWID_PERFMON -c program.c
```



Activate LIKWID
macros (C only)

Link:

```
cc -L /path/to/liblikwid program.o -o program -llikwid
```

Run:

```
likwid-perfctr -C <CPULIST> -g <GROUP> -m ./program
```



Activate
markers

MPI:

```
likwid-mpirun -np 4 -pin <PINEXPR> -g <GROUP> -m ./program
```

→ One separate block of output for every marked region

So... what should I look at first?

Focus on **resource utilization** and **instruction decomposition**!

Metrics to look at:

- Operation throughput (Flops/s)
- Overall instruction throughput (IPC,CPI)
- **Instruction breakdown**:
 - FP instructions
 - loads and stores
 - branch instructions
 - other instructions
- Instruction breakdown w.r.t. **SIMD width** (scalar, SSE, AVX, AVX512 for x86)
- **Data volumes** and **bandwidths** to main memory (GB and GB/s)
- Data volumes and bandwidth to different cache levels (GB and GB/s)

Useful diagnostic metrics are:

- Clock frequency (GHz)
- Power (W)

All the above metrics can be acquired using these performance groups:

MEM_DP, MEM_SP, BRANCH, DATA, L2, L3

Example: triangular matrix-vector multiplication

```
#define N 10000 // matrix in memory
#define ROUNDS 10
// Initialization
fillMatrix(mat, N*N, M_PI);
fillMatrix(bvec, N, M_PI);

// Calculation loop
#pragma omp parallel
{
    for (int k = 0; k < ROUNDS; k++) {
        #pragma omp for private(current,j)
        for (int i = 0; i < N; i++) {
            current = 0;
            for (int j = i; j < N; j++)
                current += mat[(i*N)+j] * bvec[j];
            cvec[i] = current;
        }
        while (cvec[N>>1] < 0) {dummy();break;}
    }
}
```



Prevent smart compilers from eliminating benchmark if `cvec` not used afterwards

Example: triangular matrix-vector multiplication

```
#include <likwid-marker.h>
[...] // defines, fillMatrix, init data
LIKWID_MARKER_INIT;
#pragma omp parallel
{
    for (int k = 0; k < ROUNDS; k++) {
        LIKWID_MARKER_START("Compute");
        #pragma omp for private(current,j)
        for (int i = 0; i < N; i++) {
            current = 0;
            for (int j = i; j < N; j++)
                current += mat[(i*N)+j] * bvec[j];
            cvec[i] = current;
        }
        LIKWID_MARKER_STOP("Compute");
        while (cvec[N>>1] < 0) {dummy();break;}
    }
}
LIKWID_MARKER_CLOSE;
```



Example: triangular matrix-vector multiplication

```
$ likwid-perfctr -C 0,1,2 -g L2 -m ./a.out
```

```
-----  
CPU type: Intel Icelake SP processor  
CPU clock: 2.39 GHz  
-----
```

```
<<<< PROGRAM OUTPUT >>>>
```

```
Region Compute, Group 1: L2
```

```
+-----+-----+-----+-----+  
|   Region Info   | HWThread 0 | HWThread 1 | HWThread 2 |  
+-----+-----+-----+-----+  
| RDTSC Runtime [s] |   0.198263 |   0.198364 |   0.198246 |  
|   call count     |          10 |          10 |          10 |  
+-----+-----+-----+-----+
```


```
+-----+-----+-----+-----+-----+  
|           Event           | Counter | HWThread 0 | HWThread 1 | HWThread 2 |  
+-----+-----+-----+-----+-----+  
| INSTR_RETIRED_ANY        | FIXC0   | 194399400 | 269695800 | 341470000 |  
| CPU_CLK_UNHALTED_CORE    | FIXC1   | 458193600 | 464605300 | 433236300 |  
| CPU_CLK_UNHALTED_REF     | FIXC2   | 473442400 | 469863600 | 465054300 |  
| TOPDOWN_SLOTS            | FIXC3   | 229096800 | 232302600 | 216618100 |  
| L1D_REPLACEMENT          | PMC0    | 69660770 | 41754150 | 7610321 |  
| L2_TRANS_L1D_WB          | PMC1    | 43768 | 263047 | 442018 |  
| ICACHE_64B_IFTAG_MISS    | PMC2    | 9698 | 11399 | 11571 |  
+-----+-----+-----+-----+-----+
```

???

Example: triangular matrix-vector multiplication

Retired instructions are misleading!

Waiting in implicit OpenMP barrier executes many instructions




We need to measure actual work (or use a tool that can separate user from runtime lib instructions)

Example: triangular matrix-vector multiplication

Floating-point instructions reliable \leftrightarrow useful work metric

Caveats:

- FP instruction counters from SandyBridge to Haswell are only qualitatively correct
- Masked SIMD lanes cannot be counted directly on x86




Example: triangular matrix-vector multiplication

Changing OMP schedule to **static** with **chunk size 16** → smaller work packages per thread


No imbalance anymore!

Is it also faster?

Triangular DMVM on Intel IcelakeSP,
Event FP_ARITH_INST_RETIRED_256B_PACKED_DOUBLE,
OMP schedule(static,16)



Triangular DMVM on Intel IcelakeSP,
Double-precision MFLOPS/s,
OMP schedule(static) vs schedule(static,16)



Summary of hardware performance monitoring

- Useful **only if you know what you are looking for**
 - Hardware event counting bears the potential of acquiring massive amounts of data for nothing!
- **Resource-based metrics** are most useful
 - Cache lines transferred, work executed, loads/stores, cycles
 - Instructions, CPI, cache misses may be misleading
- Caveat: **Processor work \neq user work**
 - Waiting time in libraries (OpenMP, MPI) may cause lots of instructions
 - → distorted application characteristic
 - Advanced tools can discern user instructions from runtime library instructions
- Another very useful application of PM: **validating performance models!**
 - Roofline is data centric → measure data volume through memory hierarchy