



Neural Network Tricks

Technische Universität Berlin - Machine Learning Group



Agenda

Hyperparameters and Optimization

Controlling Overfitting

Tips, Tricks and Misc

Summary






Motivation

- Large number of parameters \rightarrow overfitting
- Non-convex, high-dimensional optimization with SGD \rightarrow local optima
- Learning from small datasets
- **Important:** Generalization error





Agenda

Hyperparameters and Optimization

Controlling Overfitting

Tips, Tricks and Misc

Summary





Training Curve

- Visualize evolution of learning
- Performance metrics: loss, accuracy, precision, recall etc.
- Training and **validation** performance




Figure: Typical SGD training curve for neural networks.





Learning Rate

- Far too high: divergence
- Too high: Fast initial improvement but convergence to bad optimum
- Appropriate: typically exponentially shaped
- Too low: Linear improvement
- Search in logarithmic space
- Typical value: 10^{-2}





Learning Rate Scheduling

- High initial learning rate
- Decrease to find minimum more accurately
- Variants:
 - Step
 - Exponential
 - $1/t$
 - Reduction on plateau





Batch Size

- Control “stochasticity” of SGD \rightarrow “noise” in training curve
- Influence mainly on training speed, not test performance
- Typical value: 64





Parameter Initialization

- Break symmetry by random initialization
- Size of weights should depend on number of input connections
- Method by Glorot and Bengio 2010 (“Xavier initialization”)





Momentum

- “Smooth out” oscillations in training curve
- Moving average of past gradient values
- Typical value: $\mu = 0.9$ (can be lower in the beginning)

$$\Lambda_{\text{mom}}^{t+1} = \mu \Lambda_{\text{mom}}^t - \lambda \cdot \nabla E(\Theta^t) , \quad (1)$$

where μ is the momentum parameter, Λ is the parameter update, Θ is the parameter vector and λ is the learning rate.





Advanced Update Functions

Automatic adaptation of learning rate

- AdaGrad: Duchi, Hazan, and Singer 2011
- AdaDelta: Zeiler 2012
- Adam: Kingma and Ba 2014
- RMSprop: Tieleman and Hinton 2012
- Nesterov's Accelerated Gradient (NAG): Nesterov 1983





Advanced Update Functions

Automatic adaptation of learning rate

- AdaGrad: Duchi, Hazan, and Singer 2011
- AdaDelta: Zeiler 2012
- **Adam**: Kingma and Ba 2014
- RMSprop: Tieleman and Hinton 2012
- Nesterov's Accelerated Gradient (NAG): Nesterov 1983





Hyperparameter Search

- Grid search
- Random search (Bergstra and Bengio 2012)
- Bayesian hyperparameter optimization





Agenda

Hyperparameters and Optimization

Controlling Overfitting

Tips, Tricks and Misc

Summary





Early Stopping

- Select model with best validation performance
- Stop training if validation performance has not improved for some epochs (“patience”)




Figure: Early stopping.





Early Stopping

- Select model with best validation performance
- Stop training if validation performance has not improved for some epochs (“patience”)




Figure: Early stopping.





Dropout

- Co-adapting neurons: “errors” of certain neurons may be compensated by others
- Dropout neurons randomly during training (Srivastava et al. 2014)
- Ensemble of pruned networks
- Averaging in test phase
- Takes longer to train (2-3 times)
- Dropout ratio ρ : probability that neuron is kept in SGD iteration
- Typical value: $\rho = 0.5$





Weight Regularization

- Regularize length of parameter vector
- L_1 and L_2
- “Weight decay”: regularization strength
- Typical value: 10^{-3}

$$J = E(\Theta) + \omega \|\Theta\|_2^2 \rightarrow \min, \quad (2)$$

where J is the objective function, E is the cost function, ω is the weight decay and Θ are the parameters of the model.





Data Augmentation

- More data always reduces overfitting
- Generate synthetic data by applying transformations, e. g. for images:
 - random crops
 - mirroring
 - rotation
 - elastic deformations
 - color augmentation





Transfer Learning

- Extract knowledge from *source task* and apply to *target task*
- Pretraining and finetuning → pretrained parameters as initialization
- Unsupervised and supervised
- Narrow down exploration space
- Acts as regularizer





Agenda

Hyperparameters and Optimization

Controlling Overfitting

Tips, Tricks and Misc

Summary





Standardization

- Features should all be in same scale and centered around zero
- Express features in terms of standard deviations from the mean

$$x' = \frac{x - \mu}{\sigma} \quad (3)$$

- Images: μ and σ can be computed per:
 - feature \rightarrow “mean image”
 - channel
 - sample
- Standardize validation and test set with training statistics





Batch Normalization

- Standardize input features in each layer by mini-batch statistics
- Enables higher learning rates
- Proposed by Ioffe and Szegedy 2015

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.





Ensembles

- Boost performance by averaging predictions of several models
- Average predictor's variance (bias-variance trade-off)
- Models should be as diverse as possible, e. g. :
 - Different architectures
 - Different initializations (→ pretraining)
 - Different optimization
 - Best models of cross validation
 - Models that are not neural networks, e. g. SVMs, Random Forests





Agenda

Hyperparameters and Optimization

Controlling Overfitting

Tips, Tricks and Misc

Summary





Recommendations

- Most important hyperparameters:
 - Initial learning rate
 - Learning rate schedule (step?)
 - Regularization strength (dropout and weight decay)
 - Batch size
- Random search better than grid search
- Overfit subset of data
- SGD + momentum or Adam
- Standardization is crucial!
- One validation fold better than k-fold cross validation for big data
- Get more data
- Get more GPUs





References I

- Bengio, Yoshua (2012). “Practical recommendations for gradient-based training of deep architectures”. In: **Neural Networks: Tricks of the Trade**. Springer, pp. 437–478.
- Bergstra, James and Yoshua Bengio (2012). “Random search for hyper-parameter optimization”. In: **Journal of Machine Learning Research** 13.Feb, pp. 281–305.
- Duchi, John, Elad Hazan, and Yoram Singer (July 2011). “Adaptive subgradient methods for online learning and stochastic optimization”. In: **Journal of Machine Learning Research** 12, pp. 2121–2159.
- Erhan, Dumitru et al. (Feb. 2010). “Why does unsupervised pre-training help deep learning?” In: **Journal of Machine Learning Research** 11, pp. 625–660.
- Glorot, Xavier and Yoshua Bengio (2010). “Understanding the difficulty of training deep feedforward neural networks”. In: **Proceedings of the thirteenth international conference on artificial intelligence and statistics**, pp. 249–256.





References II

- Ioffe, Sergey and Christian Szegedy (2015). “Batch normalization: Accelerating deep network training by reducing internal covariate shift”. In: **arXiv preprint arXiv:1502.03167**.
- Karpathy, Andrej (n.d.). **CS231n Convolutional Neural Networks for Visual Recognition**.
<http://cs231n.github.io/neural-networks-3/>.
- Kingma, Diederik and Jimmy Ba (2014). “Adam: A method for stochastic optimization”. In: **arXiv preprint arXiv:1412.6980**.
- Nesterov, Yurii (1983). “A method of solving a convex programming problem with convergence rate $O(1/k^2)$ ”. In: **Soviet Mathematics Doklady**. Vol. 27. 2, pp. 372–376.
- Pan, Sinno Jialin and Qiang Yang (2010). “A survey on transfer learning”. In: **IEEE Transactions on knowledge and data engineering** 22.10, pp. 1345–1359.
- Srivastava, Nitish et al. (2014). “Dropout: a simple way to prevent neural networks from overfitting.”. In: **Journal of Machine Learning Research** 15.1, pp. 1929–1958.





References III

- Tieleman, Tijmen and Geoffrey Hinton (2012). “Divide the gradient by a running average of its recent magnitude”. In: **COURSERA: Neural Networks for Machine Learning 4.2**.
- Zeiler, Matthew D (2012). “ADADELTA: an adaptive learning rate method”. In: **arXiv preprint arXiv:1212.5701**.





Thank you!

