

Lecture 9

**Recurrent Neural Networks**

## Stateless vs. Stateful Models

- ▶ Characterization
- ▶ Examples

## Recurrent Neural Networks (RNNs)

- ▶ General formulation of a RNN
- ▶ Examples of practical RNNs  
(e.g. standard, bidirectional, encoder-decoder)
- ▶ Choosing the initial state

## The Difficulty of Training RNNs

- ▶ The vanishing/exploding gradient problem

## LSTM Architecture for RNNs

## Applications of RNNs

Part 1


## **Stateless vs. Stateful Models**

# Stateless vs. Stateful Models

---


## Stateless Predictor

- ▶ The prediction is simply a function of the data given as input.
- ▶ The data given as input could be e.g. a simple vector of measurement, or a sequence of such vectors (a time series).



## Stateful Predictor

- ▶ The prediction is a function that produces a prediction from the input and the current state of the system. The function also outputs the future state of the system.



# Stateless vs. Stateful Models

---

## Example: Stateless model for moving average

- ▶ Equation:

$$y_t = \alpha \cdot x_t + \beta x_{t-1} + \gamma x_{t-2}$$

- ▶ This model can be interpreted as a sliding window through the input sequence, and has a finite horizon.
- ▶ Assuming an input time series  $(x_1, x_2, \dots, x_T)$ , values  $y_3, y_4, \dots$  can be predicted directly.
- ▶ This equation can be modeled using a Convolutional Neural Network (CNN).

## Example: Stateful model for moving average

- ▶ Equation:

$$\begin{bmatrix} y_t \\ h_t \end{bmatrix} = \begin{bmatrix} h_t \\ \gamma h_{t-1} + (1 - \gamma)x_t \end{bmatrix}$$

- ▶ This model has an infinite horizon.
- ▶ Assuming an input time series  $(x_1, x_2, \dots, x_T)$  one needs to specify an initial state  $h_0$  to compute any of the predicted values  $y_t$ .


# Stateful Models

## Stateful Models

- ▶ Can be useful when we are dealing with sequential data like natural language, audio, stock prices, and etc.

## Questions

- ▶ How can we build a network that can solve the second equation of the previous slide?



Part 2

## **Recurrent Neural Networks**

# Towards a General Formulation: RNNs

---

- ▶ The model studied above can be generalized by the equation:

$$\begin{bmatrix} \mathbf{y}_t \\ \mathbf{h}_t \end{bmatrix} = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \cdot \begin{bmatrix} \mathbf{x}_t \\ \mathbf{h}_{t-1} \end{bmatrix}.$$

The matrices  $A, B, C, D$  can be learned from the data, e.g. to minimize the divergence between the output time series  $\mathbf{y}$  and some ground-truth time series  $\mathbf{t}$ .

- ▶ The model above can be further generalized to:

$$\begin{bmatrix} \mathbf{y}_t \\ \mathbf{h}_t \end{bmatrix} = f_{\theta} \left( \begin{bmatrix} \mathbf{x}_t \\ \mathbf{h}_{t-1} \end{bmatrix} \right)$$

where  $f_{\theta}$  can be any function, e.g. a neural network, with a set of parameters  $\theta$  to be learned.




# RNN Visualization

The Vanilla RNN defined by the equations:

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$

$$y_t = W_{hy}^T h_t$$

can be visualized as:



## Observation:

- ▶ A RNN can be seen as a big neural network composed of a large number of sub neural networks with shared parameters. The whole architecture can be trained via backprop.
- ▶ The function  $f_\theta$  is composed of multiple times. If  $f_\theta$  is a neural network of depth  $L$ , the RNN becomes a network of depth  $L \cdot T$ .

# Sequence-to-Sequence (Seq2Seq)

---

- Converting an input sequence of tokens into an output sequence of tokens.

## Machine Translation

### Example 1:

Input sequence (source language): "Ich spreche Deutsch."

Output sequence (target language): "I speak German."


### Example 2:

Input sequence (source language): "Ich gehe morgen ins Kino."

Output sequence (target language): "I am going to the cinema tomorrow."


# RNN Architectures for Sequence-to-Sequence

## Standard (unidirectional) RNNs:




- ▶ Generate the output at the same time as the input is received → enable a strong coupling between the two sequences.
- ▶ Cannot use information about later time steps when generating the output sequence (problem for e.g. translation).

## Bidirectional RNNs




- ▶ Add a RNN in reverse direction in order to incorporate information from future values in the sequence.

## Encoder-Decoder RNNs:



- ▶ Instead of generating the output sequence at the same time as we process the input sequence, first create a global representation of the input sequence  $s$ , and then, generate the output sequence from  $s$ .
- ▶ This ability to read through the whole sequence before generating is useful for tasks such as machine translation.

# The Problem of Initial States



## Problem:

- ▶ Unlike the input data, the RNN's initial state (at time  $t = 0$ ) is not given and must be initialized to some value.


## Possible approaches:

- ▶ Set it to some arbitrary value (e.g.  $h_0 = \mathbf{0}$ ).
- ▶ Set it at random (the RNN will then learn to desensitize itself to the initial state).
- ▶ Use one of the two approaches above **and** simulate the RNN for a few time steps in order to generate an initial state that is more plausible.

## Part 3

# **Difficulty of RNN Training**

# RNN Optimization: Pathological Gradients




The objective to optimize for a RNN is typically expressed as:

$$\mathcal{E} = \ell(y_1, t_1) + \dots + \ell(y_T, t_T)$$

The gradient of the objective w.r.t. the parameter vector  $\theta$  can be expressed via the chain rule:

$$\frac{\partial \mathcal{E}}{\partial \theta} = \sum_{t=1}^T \frac{\partial \mathcal{E}}{\partial y_t} \cdot \left( \frac{\partial^+ y_t}{\partial \theta} + \frac{\partial y_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial \theta} \right)$$
$$\frac{\partial h_{t-1}}{\partial \theta} = \sum_{s=2}^{t-1} \underbrace{\left( \prod_{i=s}^{t-1} \frac{\partial h_i}{\partial h_{i-1}} \right)}_{P_{s,t}} \frac{\partial^+ h_{s-1}}{\partial \theta}$$

# RNN Optimization: Pathological Gradients



## Observation:

- ▶ In the previous slide, we could express the error gradient  $\partial \mathcal{E} / \partial \theta$  as a sum over indices  $t = 1 \dots T$ , and  $s = 2 \dots t - 1$ , where each summand contains a product structure of the type.

$$P_{s,t} = \left( \prod_{i=s}^{t-1} \frac{\partial h_i}{\partial h_{i-1}} \right)$$

- ▶ On one extreme, the summand corresponding to indices  $s = 2$  and  $t = T$  features a very large product structure of  $T - 2$  terms.
- ▶ On the other extreme, for summands where  $s = t - 1$  the product structure totally vanishes (and just becomes an identity matrix  $I$ ).



# RNN Optimization: Pathological Gradients

## Analysis for the Linear Model:

- ▶ Recall that the linear model is given by the equations:

$$\begin{bmatrix} y_i \\ h_i \end{bmatrix} = \begin{bmatrix} A & B \\ C & D \end{bmatrix} \cdot \begin{bmatrix} x_i \\ h_{i-1} \end{bmatrix}$$

- ▶ For such a model, the matrix  $P_{s,t}$  can be computed in closed form:

$$P_{s,t} = \left( \prod_{i=s}^{t-1} \frac{\partial h_i}{\partial h_{i-1}} \right) = D^{t-s}$$

hence,  $P_{2,T} = D^{T-2}$ .

## Eigenvalue Decomposition

If  $D$  is diagonalizable, the matrix can be rewritten as  $D = Q\Lambda Q^{-1}$  with  $\Lambda$  containing the eigenvalues of  $D$ , then

$$D^2 = Q\Lambda \underbrace{Q^{-1}Q}_I \Lambda Q^{-1} = Q\Lambda^2 Q^{-1}$$

and after a few steps,  $D^{T-2} = Q\Lambda^{T-2}Q^{-1}$ .

# RNN Optimization: Pathological Gradients

---

## Two cases for the Linear RNNs:

$\max_k \lambda_k > 1$  The norm of the matrix  $D^{T-2}$  will keep increasing as  $T$  becomes large  $\rightarrow$  gradients tend to explode.

$\max_k \lambda_k < 1$  The norm of the matrix  $D^{T-2}$  will keep decreasing as  $T$  becomes large  $\rightarrow$  gradients tend to vanish.

## Possible Solutions:


- ▶ Utilizing gradient clipping helps mitigating the issue of exploding gradients.
- ▶ Choosing a particular class of functions for the RNN that is shown to be more robust to the vanishing/exploding gradient problem.

Part 4

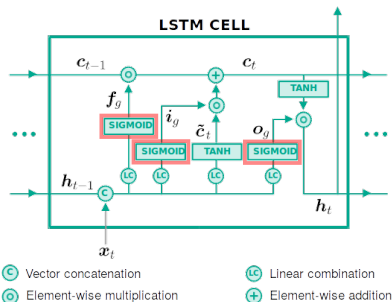
## **Long Short-Term Memory**

## Long Short-Term Memory

- ▶ The LSTM is an enhanced RNN architecture where the building blocks (cells) are equipped with special functions to stabilize learning, particularly by alleviating the issue of vanishing/exploding gradients.
- ▶ The LSTM cell, in comparison to a standard RNN cell, has an additional (more stable) state  $\mathbf{c}_t$ , that is only accessed through gate functions.



# Long Short-Term Memory



$$f_g = \sigma(W_f^T x_t + U_f^T h_{t-1} + b_f)$$

$$i_g = \sigma(W_i^T x_t + U_i^T h_{t-1} + b_i)$$

$$o_g = \sigma(W_o^T x_t + U_o^T h_{t-1} + b_o)$$

$$c'_t = \tanh(W_c^T x_t + U_c^T h_{t-1} + b_c)$$

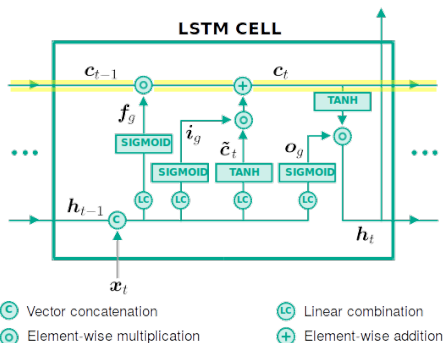
$$c_t = f_g \odot c_{t-1} + i_g \odot c'_t$$

$$h_t = o_g \odot \tanh(c_t)$$

## Observation:

- ▶ The state  $c$  is only accessed through three gates (a gate is a multiplication by a sigmoid). The 'forget gate'  $f_g$  performs an 'erase' operation. The 'input gate'  $i_g$  performs a 'write' operation. The 'output gate'  $o_g$  performs a 'read' operation.


# Long Short-Term Memory



## Observation (2):

- ▶ The state  $c$  stays stable over time (it is only erased or updated when the input gate is open), and there are no weight matrices or nonlinearities transforming  $c$  over different time steps, i.e. by default it stays constant.

# Long Short-Term Memory



## Observation (3):

- ▶ The gradient flows well and predictably along the path  $c_{t-1}, c_t, \dots$ . In particular, the addition operation does not change the gradient. The gradient can then only be dampened by the forget gate, and *never* amplified.

Part 5

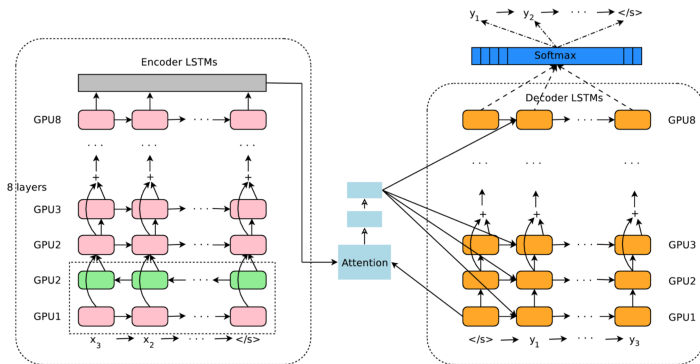
## **RNN Applications**



# RNNs for Machine Translation

## Google Neural Machine Translation:

- ▶ Encoder-Decoder architecture with input word vectors in the source language, and output in the target language.
- ▶ Stack of LSTMs with residual connections through the stack for better gradient flow. First layer is bidirectional.
- ▶ Many more details (attention mechanisms, few-shot learning procedure, etc.)




# RNNs for Modeling Motion

---


## Idea:

- ▶ Learn a recurrent neural network model of motion (e.g. of a salamander) from observed behavior.
- ▶ The motion can then be steered by forcing certain neurons or input of the RNN to take specific values.
- ▶ The model can be analyzed for insights into the mechanisms of locomotion.



# RNNs for Image Captioning

- ▶ A pre-trained CNN can be used to extract high-level features from the input image and produce an image representation.
- ▶ The image representation is passed to the RNN as its initial state.



## Summary

# Summary







---

- ▶ Recurrent neural networks (RNNs) are a special type of neural networks where the internal representation depends both on the input and on the neural network's state.
- ▶ RNNs are therefore time-dependent. This makes them natural architectures for modeling processes over time such as the evolution of dynamical systems or more generally sequential data.
- ▶ RNNs can be unfolded in time, resulting in deep neural networks with a number of layers proportional to the number of time steps, and shared parameters between the multiple layers.
- ▶ In practice, RNNs are hard to train due to the vanishing/exploding gradient problem. A powerful extension of RNNs that exhibits higher stability is the LSTM.

## References

# References

---

-  Bengio, Y., P. Simard, and P. Frasconi (1994). "Learning long-term dependencies with gradient descent is difficult". In: *IEEE Transactions on Neural Networks*.
-  Hochreiter, Sepp and Jürgen Schmidhuber (1997). "Long short-term memory". In: *Neural computation*.
-  Ijspeert, A.J., Alessandro Crespi, and Jean-Marie Cabelguen (2005). "Simulation and Robotics Studies of Salamander Locomotion: Applying Neurobiological Principles to the Control of Locomotion in Robots". In: *Neuroinformatics*.
-  Karpathy, Andrej and Li Fei-Fei (2015). "Deep visual-semantic alignments for generating image descriptions". In: *Proceedings of the IEEE conference on computer vision and pattern recognition*.
-  Pascanu, Razvan, Tomas Mikolov, and Yoshua Bengio (2013). "On the difficulty of training recurrent neural networks". In: *International Conference on Machine Learning*.
-  Wu, Yonghui et al. (2016). "Google's neural machine translation system: Bridging the gap between human and machine translation". In: *arXiv preprint arXiv:1609.08144*.