

GPU Programming II

Andrey Alekseenko

KTH Royal Institute of Technology & SciLifeLab

Refresher: Getting started on Dardel (lazy)



```
$ ssh abcd@dardel.pdc.kth.se
$ ml PDC/22.06 adaptivecpp/23.10.0-cpeGNU-22.06-rocm-5.3.3-11vm
$ export SLURM_ACCOUNT=edu23.aqti SLURM_TIMELIMIT=00:05:00
$ export SLURM_PARTITION=gpu SLURM_RESERVATION=lab-2023-12-05
$ srun acpp-info -l
=====Backend information=====
Loaded backend 0: OpenMP
  Found device: hipSYCL OpenMP host device
Loaded backend 1: HIP
Found device: AMD MI250X
Found device: AMD MI250X
Found device: AMD MI250X
Found device: AMD MI250X
```

Refresher: Getting started on TCBLab (interactive mode)

```
$ ssh wsXX@login.tcblab.org
$ salloc # Run only once; if disconnected, ssh to the same node
salloc: Nodes gpuYY are ready for job
$ ssh gpuYY
$ module load adaptivecpp/23.10.0-clang16-cuda12.1
$ acpp-info -l
=====Backend information=====
Loaded backend 0: CUDA
Found device: NVIDIA RTX A5000
Loaded backend 1: OpenMP
Found device: hipSYCL OpenMP host device
```

Exercise: Matrix transpose

- Build and run `transpose_matrix_v0.cpp`
- For now, it just copies the matrix
 - The self-check will fail!
- Look at the source code
 - See the new constructs we learned today
- Modify the code to transpose the matrix
- Use local memory to achieve coalesced memory access
- Solutions:
 - Naive: `transpose_matrix_v1.cpp`
 - Optimized: `transpose_matrix_v2.cpp`





Exercise: Matrix transpose

```
auto copyKernel(const float *in, float *out, int width, int height) {  
    return [=](sycl::nd_item<2> item) {  
        int x_index = item.get_global_id(1);  
        int y_index = item.get_global_id(0);  
        int index = y_index * width + x_index;  
        out[index] = in[index];  
    };  
}
```

```
int index  = Y * width  + X // :)  
int indexT = X * height + Y // :(
```

Exercise: Matrix transpose



Exercise: Matrix transpose

```
int index  = Y * width  + X // :)
int indexT = X * height + Y // :(
```


```
auto transposeKernel(sycl::handler &cgh,
                    const float *in, float *out,
                    int width, int height) {
    sycl::local_accessor<float, 1> tile{{tile_dim * tile_dim}, cgh};
    return [=](sycl::nd_item<2> item) {
        int x_tile_index = item.get_group(1) * tile_dim;
        int y_tile_index = item.get_group(0) * tile_dim;
        int x_local_index = item.get_local_id(1);
        int y_local_index = item.get_local_id(0);

        tile[...] = in[...]; // TODO: Coalesced read from in
        item.barrier();
        out[...] = tile[...]; // TODO: Coalesced write to out
    };
}
```


Stencil computation

- Heat flow in a 2D area
- Continuous eqn. is discretized in space and time
 - We'll ignore the finer aspects of it
- Higher spatial resolution necessitates smaller time steps

$$\frac{\partial u}{\partial t} = \alpha \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right)$$



Stencil computation



Code

- `/cfs/klemming/home/a/andreya1/Public/stencil/` (Dardel)
- `/mnt/cephfs/home/aqtiivate-ws/stencil/` (TCBLab)
 - `openmp/`
 - On Dardel, `module load PrgEnv-amd/8.3.3` first
 - `core.cpp`, `main.cpp`: OpenMP on CPU
 - Use `OMP_NUM_THREADS=4` env. variable to set number of CPU threads to 4
 - `core-naive.cpp`: OpenMP on GPU, naive version
 - `core-data.cpp`, `main-data.cpp`: OpenMP on GPU, fast version
 - `sycl/`
 - `core-naive.cpp`, `main-naive.cpp`: SYCL on GPU, naive version
 - `core-data.cpp`, `main-data.cpp`: SYCL on GPU, fast version
 - (you will get runtime warnings, they are harmless)