

# GPU Programming II

Andrey Alekseenko

KTH Royal Institute of Technology & SciLifeLab

# Quiz time!

Which statement is false?

- A. Host (CPU) schedules tasks on the device (GPU).
- B. Host and device cannot access each others memory.
- C. CPUs have more control logic per core than GPUs.
- D. GPUs tend to offer more FLOPs than CPUs (per Watt/\$).

# Quiz time!

Which statement is false?

- A. Host (CPU) schedules tasks on the device (GPU).
- B. Host and device cannot access each others memory.
- C. CPUs have more control logic per core than GPUs.
- D. GPUs tend to offer more FLOPs than CPUs (per Watt/\$).

# Quiz time!

What is the significance of over-subscribing the GPU?

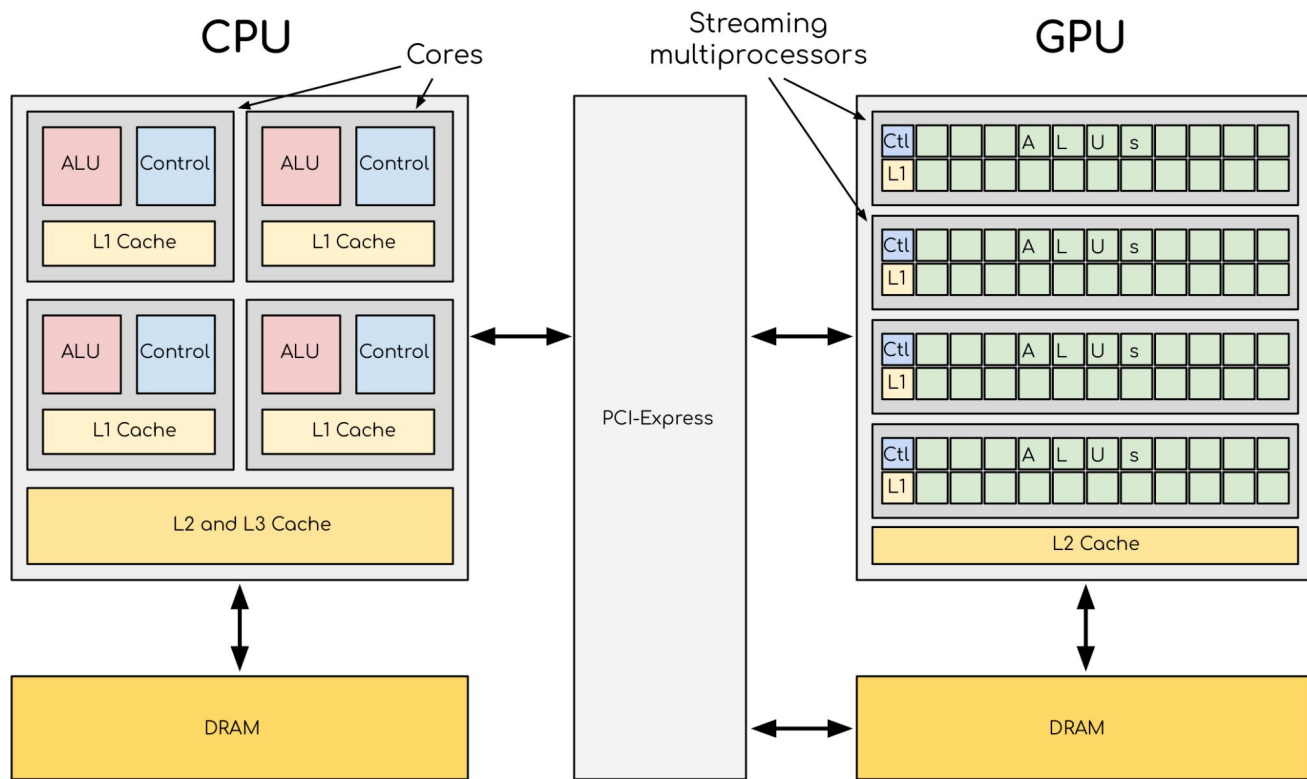
- A. It reduces the overall performance of the GPU.
- B. It helps to hide latencies and ensure high occupancy of the GPU.
- C. It leads to a memory overflow in the GPU.
- D. It ensures that there are more cores than work-groups present on the device.

# Quiz time!

What is the significance of over-subscribing the GPU?

- A. It reduces the overall performance of the GPU.
- B. It helps to hide latencies and ensure high occupancy of the GPU.
- C. It leads to a memory overflow in the GPU.
- D. It ensures that there are more cores than work-groups present on the device.



# Refresher: GPU architecture



Generated with AI / DALL-E 3 /  
Microsoft Image Creator

# Refresher: Data parallelism

- **Work-item** is a single thread of execution; it has its own **private** memory
- Work-items in a **work-group** can share data via **local** memory
- All work-items in all work-groups have to **global** memory
- All work-items run the same **kernel**



# Refresher: SYCL vector add

```
sycl::queue q{{sycl::property::queue::in_order()}};

float *Ad = sycl::malloc_device<float>(N, q); // Allocate the arrays on GPU
q.copy<float>(Ah.data(), Ad, N); // Copy the input data from host to the device

sycl::range<1> global_size(N); // Define grid dimensions
// Run our kernel
q.submit([&](sycl::handler &h) {
    h.parallel_for<class VectorAdd>(global_size, [=](sycl::item<1> threadId) {
        int tid = threadId[0]; // Get thread index
        Cd[tid] = Ad[tid] + Bd[tid]; // Do the math
    });
});

q.copy<float>(Cd, Ch.data(), N); // Copy results back to the host
// All the operations before were asynchronous!
q.wait(); // Wait for the copy to finish
sycl::free(Ad, q); // Free the GPU memory
```



# Events

- Most queue submissions return an event

```
sycl::event ev_copy = q.copy<float>(...);  
sycl::event ev_kernel = q.submit([&](...) { ... }); // Will run after copy
```

- They can be used for fine-grained synchronization

```
// Will wait for the copy, but not for the kernel  
ev_copy.wait();
```

- SYCL events are similar to OpenCL events, but different from CUDA/HIP events.

# Events

- Events can also be used for timing the kernels

```
sycl::queue q{{sycl::property::queue::in_order(),  
              sycl::property::queue::enable_profiling()}};  
  
sycl::event ev_kernel = q.submit([&](...) { ... });  
  
uint64_t kernel_start = ev_kernel  
    .get_profiling_info<sycl::info::event_profiling::command_start>();  
uint64_t kernel_end = ev_kernel  
    .get_profiling_info<sycl::info::event_profiling::command_end>();  
  
uint64_t kernel_duration_ns = kernel_end - kernel_start;
```

# Work-item indexing

- So far, we only dealt with flat, 1-dimensional indexing:

```
q.submit([&](sycl::handler &cgh) {  
    cgh.parallel_for<class Kernel>(sycl::range<1>{N},  
    [=](sycl::item<1> threadId) {  
        int index = threadId[0];  
    }  
    );  
});
```

- But we often deal with 2D or 3D data
- And what about work-groups and local memory?

# Multi-dimensional indexing

```
sycl::range<1> kernel_range{N}
q.submit([&](sycl::handler &cgh) {
    cgh.parallel_for<class Kernel1D>(
        kernel_range,
        [=](sycl::item<1> threadId) {
            int index = threadId[0];
            out[index] = in[index];
        });
});
```

```
sycl::range<2> kernel_range{W, H};
q.submit([&](sycl::handler &cgh) {
    cgh.parallel_for<class Kernel2D>(
        kernel_range,
        [=](sycl::item<2> threadId) {
            int x_id = threadId[0];
            int y_id = threadId[1];
            int index = x_id * H + y_id;
            out[index] = in[index];
        });
});
```

# Multi-dimensional indexing

- Can be up to 3D
- It is just syntactic sugar!

```
sycl::item<1> idx{x}, sycl::range<1> r{A};  
int index = idx[0] = x;
```

```
sycl::item<2> idx{x, y}, sycl::range<2> r{A, B};  
int index = idx[0] * r[1] + idx[1] = x * B + y;
```

```
sycl::item<3> idx{x, y, z}, sycl::range<3> r{A, B, C};  
int index = (idx[0] * r[1] + idx[1]) * r[2] + id[2]  
           = (x * B + y) * C + z;
```

# Multi-dimensional indexing

- Can be up to 3D
- It is just syntactic sugar!

```
sycl::item<1> idx{x}, sycl::range<1> r{A};  
int index = idx[0] = x;
```

```
sycl::item<2> idx{x, y}, sycl::range<2> r{A, B};  
int index = idx[0] * r[1] + idx[1] = x * B + y;
```

```
sycl::item<3> idx{x, y, z}, sycl::range<3> r{A, B, C};  
int index = (idx[0] * r[1] + idx[1]) * r[2] + id[2]  
            = (x * B + y) * C + z;
```

```
int index = idx.get_linear_id();
```

# Multi-dimensional indexing

- Can be up to 3D
- It is just syntactic sugar!
- Portability warning: order is different in CUDA/HIP/OpenCL!

```
// SYCL
sycl::item<2> idx{x, y}; sycl::range<2> r{A, B};
int index = idx[0] * r[1] + idx[1] = x * B + y;
```


```
// CUDA/HIP
int2 idx{x, y}; dim2 r{A, B};
int index = idx.y * r.y + idx.x = y * A + x;
```

# Work-group indexing

$$\text{global\_range} = 256 * 4096$$

$$\text{local\_range} = 256$$

$$\text{group\_range} = 4096$$



$$\text{global\_id} = \text{group\_id} * \text{local\_range} + \text{local\_id}$$

$$\text{global\_id} = 1 * 256 + 2 = 258$$



# Range vs. NDRange

```
sycl::range<1> global_range{n}; // n work-items in total
q.submit([&](sycl::handler &cgh) {
    cgh.parallel_for(global_range, [=](sycl::item<1> item) {
        int global_id = item.get_id(0); // [0; n)
    });
});
```



# Range vs. NDRange

```
sycl::range<1> global_range{n}; // n work-items in total
q.submit([&](sycl::handler &cgh) {
    cgh.parallel_for(global_range, [=](sycl::item<1> item) {
        int global_id = item.get_id(0); // [0; n)
    });
});
```

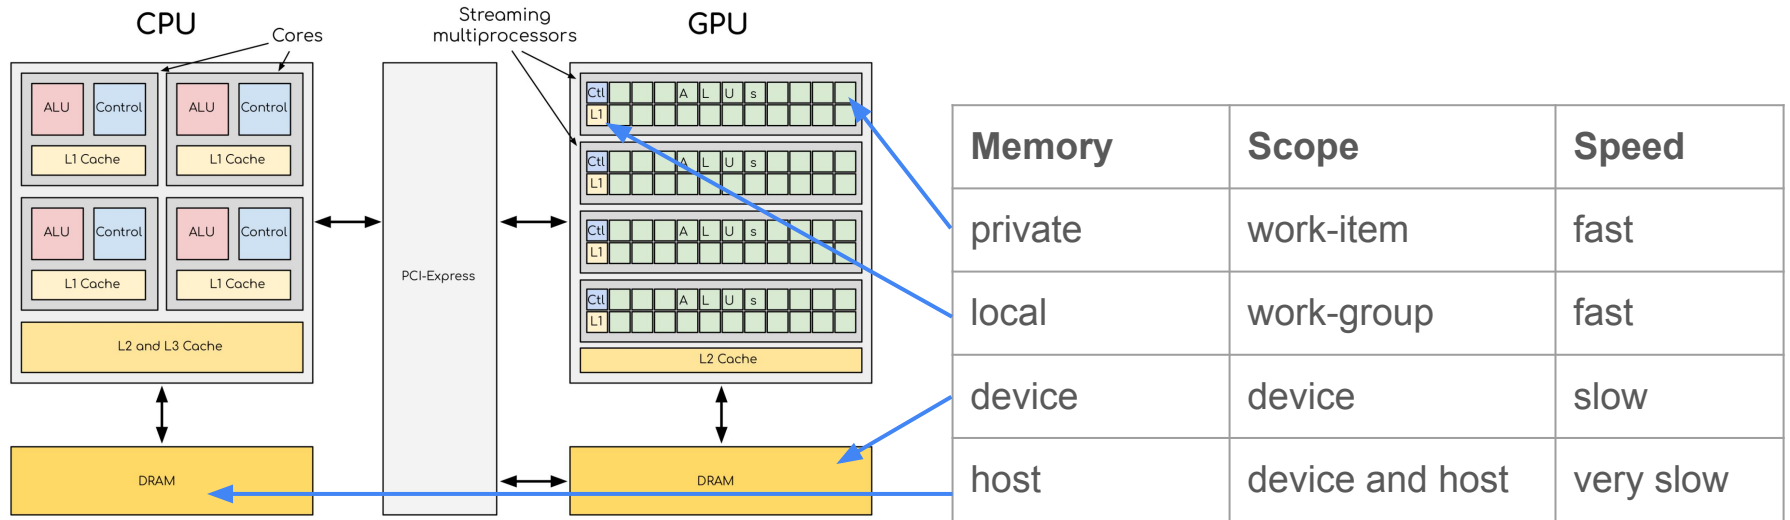
```
sycl::range<1> global_range{n}; // n work-items in total
sycl::range<1> local_range{k}; // k work-items in each work-group
sycl::nd_range<1> kernel_range{global_size, local_size};
q.submit([&](sycl::handler &cgh) {
    cgh.parallel_for(kernel_range, [=](sycl::nd_item<1> item) {
        int local_id = item.get_local_id(0); // 0..k
        int group_id = item.get_group(0); // 0..n/k
        int global_id = item.get_global_id(0); // 0..n
    });
});
```

# Why NDRange?

- Local memory
- Collective operations
- Another knob for performance tuning



# Local memory



# Local memory

SYCL	Scope	Speed
private	work-item	fast
local	work-group	fast
device	device	slow

```
q.submit([&](sycl::handler &cgh) {  
    cgh.parallel_for(kernel_range, [=](sycl::nd_item<1> item) {  
        int local_id = item.get_local_id(0); // 0..wg_size  
        int global_id = item.get_global_id(0); // 0..n  
        float x1 = input_data[global_id];  
        float x2 = input_data[global_id + wg_size - 2 * local_id];  
    });  
});
```



# Local memory

SYCL	Scope	Speed
private	work-item	fast
local	work-group	fast
device	device	slow

```
q.submit([&](sycl::handler &cgh) {  
    cgh.parallel_for(kernel_range, [=](sycl::nd_item<1> item) {  
        int local_id = item.get_local_id(0); // 0..wg_size  
        int global_id = item.get_global_id(0); // 0..n  
        float x1 = input_data[global_id];  
        float x2 = input_data[global_id + wg_size - 2 * local_id];  
    });  
});
```

# Coalesced memory access

- When adjacent threads access adjacent elements, operations are combined
- Exact rules are complicated and hardware dependent




# Local memory

SYCL	Scope	Speed
private	work-item	fast
local	work-group	fast
device	device	slow


```
q.submit([&](sycl::handler &cgh) {  
    sycl::local_accessor<float, 1> local_buffer{wg_size, cgh};  
    cgh.parallel_for(kernel_range, [=](sycl::nd_item<1> item) {  
        int local_id = item.get_local_id(0); // 0..wg_size  
        int global_id = item.get_global_id(0); // 0..n  
        local_buffer[local_id] = input_data[global_id];  
        item.barrier();  
        float x2 = local_buffer[wg_size - local_id];  
    });  
});
```




# Matrix multiplication




# Matrix multiplication




# Matrix multiplication



# Matrix multiplication




# Matrix multiplication



There are also specialized hardware units...


Just use DGEMM

# Reductions




# Reductions

```
q.submit([&](sycl::handler &cgh) {  
    sycl::local_accessor<int, 1> buf{{wg_size}, cgh};  
    cgh.parallel_for(kernel_range,  
        [=](sycl::nd_item<1> item) {  
        int id = item.get_local_id(0);  
        buf[id] = x[item.get_global_id(0)];  
        for (int s = wg_size / 2; s > 0; s /= 2) {  
            item.barrier();  
            if (id < s) {  
                buf[id] += buf[id + s];  
            }  
        }  
        if (id == 0) {  
            sum[item.get_group(0)] = buf[0];  
        }  
    });  
});
```



# Reductions

```
q.submit([&](sycl::handler &cgh) {
    sycl::local_accessor<int, 1> buf{{wg_size}, cgh};
    cgh.parallel_for(kernel_range,
        [=](sycl::nd_item<1> item) {
            int id = item.get_local_id(0);
            buf[id] = x[item.get_global_id(0)];
            for (int s = wg_size / 2; s > 0; s /= 2) {
                item.barrier();
                if (id < s) {
                    buf[id] += buf[id + s];
                }
            }
            if (id == 0) {
                sum[item.get_group(0)] = buf[0];
            }
        });
});
```



Can also use sub-groups...



# Quiz time!

Which of the following programs will **not** benefit from using local memory?

- A. Matrix transpose.
- B. Matrix multiplication.
- C. Vector addition.
- D. Vector dot product.

# Quiz time!

Which of the following programs will **not** benefit from using local memory?

- A. Matrix transpose.
- B. Matrix multiplication.
- C. Vector addition.
- D. Vector dot product.

# Dealing with large kernels

```
int do_compute() {
    sycl::queue q;
    // allocate memory
    // initialize data
    q.submit([&](sycl::handler &h) {
        h.parallel_for<class ComplexKernelA>(global_size_1,
            [=](sycl::id<1> threadId) {
                // Do lots of math
                // Like, really a lot
                // Maybe a thousand lines
            }
        );
    });
    // do things
    q.submit([&](sycl::handler &h) {
        h.parallel_for<class ComplexKernelB>(global_size_2,
            [=](sycl::id<1> threadId) {
                // Oh no, more code
                // Much more
            }
        );
    });
    // copy back
}
```

# Dealing with large kernels

```
int main() {  
    q.submit([&](sycl::handler &h) {  
        h.parallel_for<class VectorAdd>(global_size,  
            [=](sycl::id<1> threadId) {  
                int tid = threadId.get(0);  
                Ad[tid] = Ad[tid] + Bd[tid];  
            })  
    });  
}
```

```
auto kernel(float *A, float *B, int N) {  
    return [=](sycl::id<1> threadId) {  
        int tid = threadId.get(0);  
        A[tid] = A[tid] + B[tid];  
    };  
}  
  
int main() {  
    q.submit([&](sycl::handler &h) {  
        h.parallel_for<class VectorAdd>(global_size,  
            kernel(Ad, Bd, N)  
        );  
    });  
}
```