




# MPI: Part II

## AQTIVATE Training Workshop I

**Dirk Pleiter**

CST | EECS | KTH

December 2023





# Overview

Buffering and Non-Blocking Communication

Halo Exchange

Collective Communication



# Content

Buffering and Non-Blocking Communication

Halo Exchange


Collective Communication

# Communication Buffering (1/3)

- ▶ MPI communication involves multiple transactions
- ▶ Challenges:
  - ▶ Intermediate buffer space required
  - ▶ Buffered data must remain unchanged until buffer is completely read

# Communication Buffering (2/3)

## ► Design option #1: Intermediate memory buffer



### ► Advantages:


- MPI send completes after fast in-memory copy

### ► Disadvantages:

- Need additional space in memory
- Performing memory copies increases pressure on memory bus

# Communication Buffering (3/3)

## ► Design option #2: Zero-copy communication



### ► Advantages:

- No additional memory space and less memory traffic

### ► Disadvantages:

- MPI send only completes once data has been transferred over network

# MPI Buffered Mode

- ▶ User can explicitly request communication to be buffered (**buffered mode**):

```
1 int MPI_Bsend(  
2     const void* buf,           /* Pointer to send buffer */  
3     int count,                 /* Number of elements */  
4     MPI_Datatype datatype,     /* Data type */  
5     int dest,                  /* Destination rank */  
6     int tag,                   /* Communication tag */  
7     MPI_Comm comm              /* Communicator */  
8 );
```

- ▶ A buffered mode send operation can be started whether or not a matching receive has been posted

# MPI Buffered Mode: User-provided Buffer

- ▶ A user may specify a buffer to be used for buffering messages sent in buffered mode:

```
1  int MPI_Buffer_attach(  
2    void* sysbuf,    /* Pointer to user allocated system buffer */  
3    int size         /* Size of buffer in Bytes */  
4  );  
5  
6  int MPI_Buffer_detach(  
7    void* sysbuf,    /* Returned pointer to system buffer */  
8    int* size        /* Returned size of buffer */  
9  );
```



# Blocking versus Non-Blocking Communication

- ▶ Blocking
  - ▶ Function returns only after completion of the associated operation
  - ▶ Examples: `MPI_Send`, `MPI_Recv`
- ▶ Non-blocking
  - ▶ Function may return before the associated operation has completed
  - ▶ Examples:
    - ▶ Immediate send: `MPI_Isend`
    - ▶ Immediate receive: `MPI_Irecv`
  - ▶ Resources (e.g. message buffers) passed to function must not be reused until the operation has completed
    - ▶ Non-blocking functions return a handle that allow to query status of the initiated operation

# MPI Immediate Send and Receive

## ► Syntax of immediate send operation

```
1  int MPI_Isend(  
2    const void* buf,           /* Send buffer */  
3    int count,                 /* Number of elements */  
4    MPI_Datatype datatype,     /* Data type */  
5    int dest,                  /* Destination rank */  
6    int tag,                   /* Communication tag */  
7    MPI_Comm comm,             /* Communicator */  
8    MPI_Request *request       /* Request handle */  
9  );
```

## ► Syntax of immediate receive operation

```
1  int MPI_Irecv(  
2    const void* buf,           /* Pointer to receive buffer */  
3    int count,                 /* Number of elements */  
4    MPI_Datatype datatype,     /* Data type */  
5    int source,                /* Rank of source */  
6    int tag,                   /* Message tag */  
7    MPI_Comm comm,             /* Communicator */  
8    MPI_Status *status,        /* Status object */  
9    MPI_Request *request       /* Request handle */  
10 );
```

# Waiting for Completion (1/2)

- ▶ To wait for a communication identified by a given request handler to complete, the blocking function `MPI_Wait` can be used:

```
1 int MPI_Wait(  
2   MPI_Request *request, /* Request handle (in/out) */  
3   MPI_Status *status   /* Status object (out) */  
4 );
```

- ▶ On return the status object is updated

- ▶ `MPI_Test` is available to perform checks without block:

```
1 int MPI_Test(  
2   MPI_Request *request, /* Request handle (in/out) */  
3   int *flag,           /* Completion flag (out) */  
4   MPI_Status *status   /* Status object (out) */  
5 );
```

- ▶ The returned `flag` is set to true when operation completed

## Waiting for Completion (2/2)

- Use `MPI_Waitall` when waiting for a set of requests to complete:

```
1 int MPI_Waitall(  
2     int count,                /* Array size */  
3     MPI_Request request[],    /* Array of requests */  
4     MPI_Status status[]      /* Array of status objects */  
5 );
```

# MPI Modes Overview (1/2)

## ▶ **Standard mode:**

- ▶ A send operation is started whether or not a matching receive has been posted
- ▶ It may complete before a matching receive is posted

## ▶ **Synchronous mode:**

- ▶ A send operation can be started whether or not a matching receive has been posted
- ▶ The send operation will complete successfully only if a matching receive is posted

## ▶ **Ready mode:**

- ▶ A send operation may be started only if the matching receive is already posted

## ▶ **Buffered mode:**

- ▶ A send operation can be started whether or not a matching receive has been posted
- ▶ It may complete before a matching receive is posted

# MPI Modes Overview (2/2)

The following variants of MPI send are available:

Mode	Blocking variant	Non-blocking variant
Standard	<code>MPI_Send</code>	<code>MPI_Isend</code>
Synchronous	<code>MPI_SSend</code>	<code>MPI_Issend</code>
Ready	<code>MPI_RSend</code>	<code>MPI_Irsend</code>
Buffered	<code>MPI_BSend</code>	<code>MPI_Ibsend</code>

# Non-Blocking Collectives

- ▶ Collective operations are typically also available in a non-blocking variant
  - ▶ Examples
    - ▶ `MPI_Ibcast`
    - ▶ `MPI_Ireduce`
- ▶ These functions return a request handle such that `MPI_Wait` or `MPI_test` can be used to check for completion



# Content

Buffering and Non-Blocking Communication


Halo Exchange

Collective Communication



# Solving the 2D Poisson in Parallel

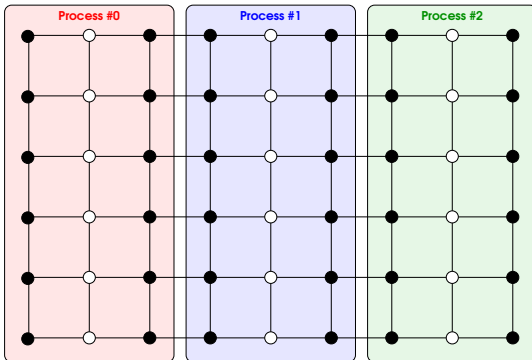
- Graphical representation of a parallel version of the discrete 2-dimensional Poisson equation:



- Local update may require data from remote process


## 2D Poisson: Halo versus Bulk

- ▶ Classification of grid points:
  - ▶ Inner points (bulk): No data from other processes needed
  - ▶ Boundary points (halo): Data from other processes needed
- ▶ Parallel update requires **halo exchange**




# 2D Poisson Equation: Data Decomposition

- Data decomposition in 1 dimension:



- Data decomposition in 2 dimensions:



# 2D Poisson Equation: Information Exchange

- ▶ Assume a square lattice of size  $L^2$  being distributed over  $P$  processes and use of double-precision numbers
- ▶ Update of each site requires 8 Flop:

$$I_{\text{fp}}(L, P) = (8 \cdot L^2 / P) \text{ Flop}$$

- ▶ In each iteration, the field  $v$  is updated and the local halo needs to be communicated
  - ▶ Data decomposition in 1 dimension:

$$I_{\text{net}}^{(1\text{d})}(L, P) = (2 \cdot L \cdot 8) \text{ Byte}$$

- ▶ Data decomposition in 2 dimensions:

$$I_{\text{net}}^{(2\text{d})}(L, P) \simeq \left( 4 \cdot (L/\sqrt{P} - 1) \cdot 8 \right) \text{ Byte}$$

# 2D Poisson Equation: Data Decomposition Analysis

- ▶ For fixed  $P$  we find

$$\frac{l_{\text{fp}}}{l_{\text{net}}} \propto L$$

- ▶ Increasing  $L$  the amount of computation relative to network communication increases

- ▶ For fixed  $L$  we find

$$\frac{l_{\text{net}}^{(1d)}}{l_{\text{fp}}} \propto P, \quad \frac{l_{\text{net}}^{(2d)}}{l_{\text{fp}}} \propto \sqrt{P}$$

- ▶ 2-dimensional data decomposition requires less data to be communicated (for sufficiently large  $L$ )
- ▶ But: In higher dimensions halo becomes fragmented in memory

## 2D Poisson Equation: Standard Send

- ▶ Consider halo update using the following pseudo-code:

```
foreach neighbour i:  
    MPI_Send(..., nb[i], ...)
```

```
foreach neighbour i:  
    MPI_Recv(..., nb[i], ...)
```

- ▶ Will this code work? Why not?

## 2D Poisson Equation: Immediate Send

- ▶ Answer: No, the send operations will start to block
- ▶ Fixed halo update using immediate send and receive operations:

```
foreach neighbour i:  
    MPI_Irecv(..., nb[i], ..., recv_req[i])  
  
foreach neighbour i:  
    MPI_Isend(..., nb[i], ..., send_req[i])  
MPI_Waitall(...)
```

# 2D Poisson Equation: Optimisations (1/2)

- ▶ Optimisation strategy #1: Overlap communication and computation
  - ▶ Steps:
    1. Initiate communication of boundary points
    2. Update interior points
    3. Wait for communications to complete
    4. Update boundary points
  - ▶ If  $\Delta t(L) = \max[\Delta t_{\text{fp}}(L), \Delta t_{\text{net}}(L)]$  and latency-bandwidth model to hold then problem becomes local performance bound for large  $L$  as

$$\Delta t \rightarrow \Delta t_{\text{fp}}(L) \text{ for } L \rightarrow \infty$$



## 2D Poisson Equation: Optimisations (2/2)

- ▶ Optimisation strategy #2: Minimise number of send-receive operations
  - ▶ Solution: Copy halo into single buffer before communication
  - ▶ Alternative solution: Use advance data types (see next section)
- ▶ Performance benefits difficult to predict as message rates improved
  - ▶ But: Higher network bandwidth demands due to additional overhead (message headers, minimal message length)



# Content


Buffering and Non-Blocking Communication

Halo Exchange

Collective Communication


# Communication Patterns: One-to-All Broadcast

- ▶ One processor has a piece of data which needs to be sent to all other processes
- ▶ Implementation assuming ring topology with  $P$  processes and process #0 being data source
- ▶ Need to perform sequence of send-receive operations until data arrives at node # $P-1$   
 $\Rightarrow$  **Complexity** =  $P - 1$  steps




# Communication Patterns: Broadcast (2)

- Implementation on 2-dimensional mesh  
 $\Rightarrow$  Complexity  $\sim 2\sqrt{P}$
- Generalization to  $d$ -dimension mesh  
 $\Rightarrow$  Complexity  $\sim d P^{1/d}$



# Communication Patterns: All-to-One Reduce

- ▶ Data on all processes is reduced to single piece of data on one processor
- ▶ Examples for reduction operators
  - ▶ Summation
  - ▶ Selection of minimum/maximum value
- ▶ Implementation assuming ring topology with  $P$  processes and process #0 being final destination:



# Communication Patterns: All-to-One Reduce (2)

- Implementation of global sum  $S$ :


$$s_k = \sum_{i=k(N/P)}^{(k+1)(N/P)-1} x_i, \quad S = \sum_{k=0}^{P-1} s_k$$

with  $s_k$  distributed over  $P$  processes connected in a ring topology:

- Process #P-1:
  - Send  $s_{P-1}$  to process #P-2
- Process #P-2:
  - Receive  $s_{P-1}$  from process #P-1
  - Compute  $s_{P-2} + s_{P-1}$
  - Send result to process #P-3
- ...
- Process #0:
  - Receive  $\sum_{k=1}^{P-1} s_k$  from process #1
  - Compute  $\sum_{k=0}^{P-1} s_k$

# Communication Patterns: Gather and Scatter

- ▶ **Scatter**: One process distributes different pieces of data to all other processes
  - ▶ Also called: one-to-all personalized communication
  - ▶ Operation fundamentally different from broadcast
- ▶ **Gather**: One process collects one piece of data from all other processes
  - ▶ Inverse of scatter operation



# MPI Collective Operations: Overview

- ▶ MPI defines 9 types of collective operations:
  - ▶ Barrier
  - ▶ Broadcast
  - ▶ Gather
  - ▶ All-Gather
  - ▶ Scatter
  - ▶ All-to-all
  - ▶ Reduce
  - ▶ All-Reduce
  - ▶ Reduce-Scatter
  - ▶ Scan
- ▶ Collective communication is over all of the processes in particular group identified by a communicator
- ▶ Standard meanwhile introduced blocking and non-blocking variants



# MPI Barrier

- ▶ Syntax:

```
int MPI_Barrier(MPI_Comm comm)
```

```
MPI_Barrier(comm, ierror)
```

```
    TYPE(MPI_Comm), INTENT(IN)      :: comm
```

```
    INTEGER, OPTIONAL, INTENT(OUT) :: ierror
```

- ▶ Function blocks the caller until all group members (identified by the communicator `comm`) have called it
  - ▶ A deadlock occurs if one group member does not call the function
- ▶ Practical advice
  - ▶ In a program with send-receive and other collective communications, an MPI barrier is rarely needed
  - ▶ MPI makes no guarantees on how long it will take other processes to leave the barrier
    - ▶ MPI barriers are not appropriate for highly accurate time measurements

# MPI Broadcast

## ► Syntax:


```

1  int MPI_Bcast(
2  void* buffer,           /* Pointer to output/input buffer */
3  int count,             /* Number of elements */
4  MPI_Datatype datatype, /* Data type */
5  int root,              /* Rank of root process */
6  MPI_Comm comm          /* Communicator */
7  );

```

- Function copies data from process identified by `root` and `comm` to all other processes of the group

- Example with 4 processes and `root=1`:




## ► Syntax:

```

1  int MPI_Gather(
2      const void* sendbuf, /* Pointer to output buffer */
3      int sendcount,      /* Number of elements */
4      MPI_Datatype sendtype, /* Data type */
5      void* recvbuf,      /* Pointer to input buffer */
6      int recvcnt,        /* Number of elements */
7      MPI_Datatype rcvtype, /* Data type */
8      int root,           /* Rank of receiving process */
9      MPI_Comm comm       /* Communicator */
10 );

```

- Function collects data on process identified by root and comm sent by all processes of the group (in rank order)



# MPI All Gather


## ► Syntax:

```

1 int MPI_Allgather(
2     const void* sendbuf, /* Pointer to send buffer */
3     int sendcount,       /* Number of elements */
4     MPI_Datatype sendtype, /* Data type */
5     void* recvbuf,       /* Pointer to recv buffer */
6     int recvcnt,         /* Number of elements */
7     MPI_Datatype recvtype, /* Data type */
8     MPI_Comm comm        /* Communicator */
9 );

```

## ► Function similar to MPI\_Gather, but now all processes receive the data




## ► Syntax:

```

1  int MPI_Scatter(
2      const void* sendbuf,    /* Pointer to output buffer */
3      int sendcount,          /* Number of elements */
4      MPI_Datatype sendtype,  /* Data type */
5      void* recvbuf,          /* Pointer to input buffer */
6      int recvcount,          /* Number of elements */
7      MPI_Datatype recvtype,  /* Data type */
8      int root,               /* Rank of sending root process */
9      MPI_Comm comm           /* Communicator */
10 );

```

## ► This function is the inverse operation to MPI\_Gather



# MPI All-to-all


## ► Syntax:

```

1 int MPI_Alltoall(
2     const void* sendbuf, /* Pointer to output buffer */
3     int sendcount,        /* Number of elements */
4     MPI_Datatype sendtype, /* Data type */
5     void* recvbuf,        /* Pointer to input buffer */
6     int recvcount,        /* Number of elements */
7     MPI_Datatype recvtype, /* Data type */
8     MPI_Comm comm         /* Communicator */
9 );

```

## ► Each process sends distinct data to each of the other processes within the group



# MPI Reduce (1/2)


## ► Syntax:

```

1 int MPI_Reduce(
2   const void* sendbuf, /* Pointer to output buffer */
3   void* recvbuf,      /* Pointer to input buffer */
4   int count,          /* Number of elements */
5   MPI_Datatype datatype, /* Data type */
6   MPI_Op op,          /* Reduction operator */
7   int root,           /* Rank of root process */
8   MPI_Comm comm       /* Communicator */
9 );

```

- Function combines the elements provided in the input buffer of each process in the group, using the operation `op`, and returns the combined value in the output buffer of the process with rank `root`



# MPI Reduce (2/2)

- ▶ MPI's predefined reduction operations:

Name	Meaning
MPI_MAX	maximum
MPI_MIN	minimum
MPI_SUM	sum
MPI_PROD	product
MPI_LAND	logical and
MPI_BAND	bit-wise and
MPI_LOR	logical or
MPI BOR	bit-wise or
MPI_LXOR	logical exclusive or (xor)
MPI_BXOR	bit-wise exclusive or (xor)
MPI_MAXLOC	max value and location
MPI_MINLOC	min value and location



# MPI Reduce Example: Global Sum

```
1  int main() {
2      const int N=1048576;
3      int isize, irank;
4
5      MPI_Init(NULL, NULL);
6      MPI_Comm_rank(MPI_COMM_WORLD, &irank);
7      MPI_Comm_size(MPI_COMM_WORLD, &isize);
8
9      int n = N / isize; /* Assume N to be multiple of isize */
10     int *v = (int *) calloc(n, sizeof(int));
11
12     for (int i = 0; i < n; i++) v[i] = irank * n + i + 1;
13
14     double s = 0; /* Local sum */
15     for (int i = 0; i < n; i++) s += v[i];
16
17     double gs = -1; /* Global sum */
18     MPI_Reduce(&s, &gs, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
19
20     if (irank == 0) printf("N=%d, gs=%.2f\n", N, gs);
21
22     free(v);
23     MPI_Finalize();
24
25     return 0;
26 }
```

# MPI All-Reduce


## ► Syntax:

```

1 int MPI_Allreduce(
2   const void* sendbuf,      /* Pointer to output buffer */
3   void* recvbuf,           /* Pointer to input buffer */
4   int count,               /* Number of elements */
5   MPI_Datatype datatype,   /* Data type */
6   MPI_Op op,               /* Reduction operator */
7   MPI_Comm comm            /* Communicator */
8 );

```

## ► Similar to MPI\_Reduce but now all processes will receive an identical copy of the result



# MPI Collectives: Caveats


- ▶ Collective operations have to be called by all processes within the given group
- ▶ Collective operations have to be issued in the same order by all participating processes
- ▶ Computations of collectives may not be deterministic
  - ▶ In exact arithmetic always the same result will be produced
  - ▶ In case of floating-point numbers rounding errors may lead to non-deterministic results
    - ▶ Re-call that in floating-point arithmetic operations may not be commutative, i.e.  $(a + b) + c \neq (c + b) + a$
  - ▶ The MPI standard does not require that the same input give the same output every time
    - ▶ Exact results are not relevant for all applications
    - ▶ No having such a requirement opens performance opportunities, e.g. by moving computations in the network such that order of the arithmetic operations will depend on network topology

# MPI Collectives: Optimisations

- ▶ For many operations, the predefined constant `MPI_IN_PLACE` can be used as send buffer argument
  - ▶ No need to allocate and manage separate send and receive buffers

# Finish with Current Networking Technology

HPE Cray Rosetta switch with 64 ports:



[Wikichip]